
Lecture 2

Parallel & Distributed Architecture

07/09/2021

IN4049TU

Parallel Computer Architectures

1. Memory organization
 - Memory hierarchy
 - Shared versus distributed memory
2. Processor/node architecture
 - Flynn's taxonomy
3. Interconnection network
 - Dynamic versus static networks


07/09/2021

IN4049TU

2

1. Classification based on Memory

1. Based on memory organization:
SM=Shared-Memory versus DM=Distributed Memory
2. Based on access time:
 - UMA =Uniform Memory Access (time);
 - NUMA=Non-uniform Memory Access (time);
 - SMP = Symmetric Multi-Processors;




07/09/2021

2. Classification based on data and control flows

Flynn's taxonomy:

1. Single Instruction Single Data (SISD) – exp. classical Von Neumann machine (sequential computer).
2. Multiple Instruction Single Data (MISD) – exp. none
3. Single Instruction Multiple Data (SIMD) – exp. GPU
4. Multiple Instruction Multiple Data (MIMD) – exp. cluster of computers




07/09/2021

IN4049TU

4

A generic parallel architecture

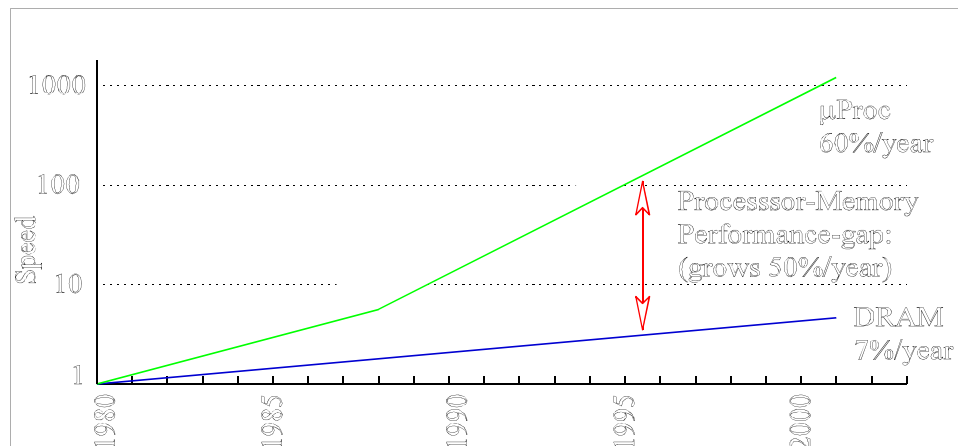


07/09/2021

IN4049TU

5

Inside one CPU: Growing gap betw. processor and memory



7-9-2021


IN4049TU

6

Memory hierarchy: Caches

In order to reduce the big delay of directly accessing the main or remote memory, it requires:

- Optimizing the data movement, maximize reuse of data already in fastest memory;
- Minimizing data movement between 'remote' memories (communication)

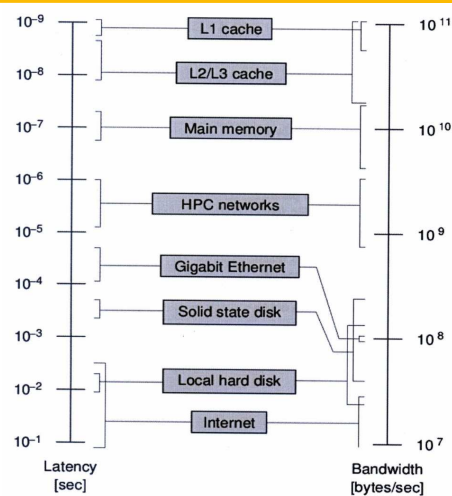


7-9-2021

IN4049TU

7

Across processors and CPUs: Speed differences in memory and networks




7-9-2021


IN4049TU

8

Extensions of Memory System (1)



(a) Shared Cache




(b) Bus-based shared memory (SMP)

07/09/2021


IN4049TU

9

Extensions of Memory System (2)



(c) Dance hall (UMA)



(d) Distributed-memory (NUMA)


07/09/2021

IN4049TU

10

Machine Model 1

- A shared memory machine
- Processors all connected to a large shared memory
- “Local” memory is not (usually) part of the hardware
 - **Symmetric Multiprocessors (SMP)**, e.g. SGI Origin
- Speed: much quicker to cache than main memory



07/09/2021

IN4049TU

11

Bus-based Shared Memory Multiprocessors

- Symmetric Multiprocessors (SMPs)
 - Symmetric access to all of main memory from any processor
- Dominate the server market
 - Building blocks for larger systems; today multi-core laptops are common
- Attractive as throughput servers and for parallel programs
 - Fine-grain resource sharing
 - Uniform access via loads/stores
 - **Automatic data movement** and coherent replication in caches
 - Useful for operating system too

07/09/2021

IN4049TU

12

Shared Memory Multiprocessors

- Normal uniprocessor mechanisms to access data through reads and writes
- Key is extension of memory hierarchy to support multiple processors (e.g., crossbar network is expensive for large number, therefore often virtual shared memory system is implemented)


07/09/2021

IN4049TU

13

Machine Model 2

- **A distributed memory machine**
 - Cluster of computers, IBM Blue Gene, Tianhe, etc.
 - Processors all connected to own memory (and caches)
 - cannot directly access another processor's memory
- **Each “node” has a network interface (NI)**
 - all communication and synchronization done through this




07/09/2021

IN4049TU

14

Machine Model 3

- A SIMD (Single Instruction Multiple Data) machine
- A large number of small processors
- A single “control processor” issues each instruction
 - each processor executes the same instruction
 - some processors may be turned off on any instruction



Earlier example machine Connection Machine (CM). Programming model is

- implemented by mapping n-fold parallelism to p processors
- mostly done in the compilers (HPF = High Performance Fortran)

07/09/2021

IN4049TU

15

Machine Model 4, Cluster of SMPs

- **Since small shared memory machines (SMP's) are the fastest commodity machine, why not build a larger machine by connecting many of them with a network?**
- Shared memory within one SMP
- Message passing outside
- ASCI Red (Intel), Blue Gene (IBM), ...
- Programming model?
 - Treat machine as “flat”, always use message passing, even within SMP (simple, but ignore important part of memory hierarchy)
 - Expose two layers: shared memory and message passing (higher performance, e.g., mixed MPI&OpenMP, but ugly to program)

07/09/2021

IN4049TU

16

Shared Memory Systems

- Shared cache
- Shared memory
- Cache coherence problem

- Emphasis is not on network, but on memory hierarchy.


07/09/2021

IN4049TU

17

Basic Shared Memory Architecture

- Processors all connected to a large shared memory
- **Local** caches for each processor
- speed: much quicker to cache than main memory




- Simplest to program, but hard to build with many processors

07/09/2021

IN4049TU

18

Limits of using Bus as Network



Assume a 1000 MB/s bus

500 MIPS processor w/o cache

=> 2000 MB/s instr BW per processor

=> 660 MB/s data BW if 33% load-store

(assuming 4 bytes per instruction/word)

Suppose 98% instr. hit rate and 95% data hit rate (assume 16 bytes block=cache line)

=> 160 MB/s instr. BW per processor

=> 132 MB/s data BW per processor

=> 292 MB/s combined BW

∴ 4 processors will saturate the bus!

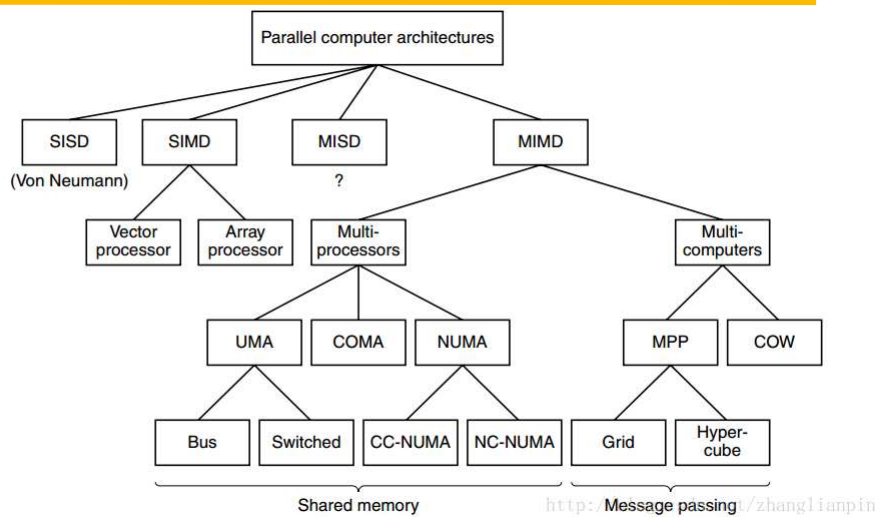
→ Bus only useful in small systems

07/09/2021

IN4049TU

19

Summary: Parallel Computer Architectures



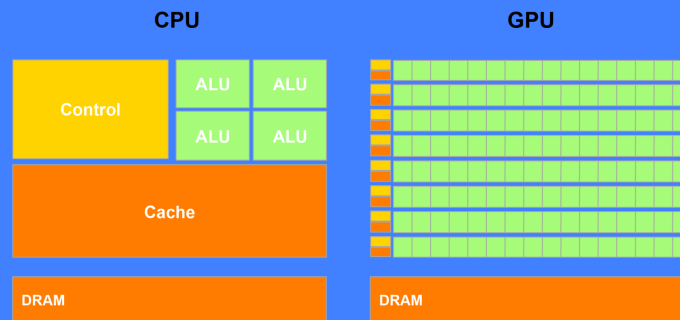
07/09/2021

IN4049TU

20

Accelerators (for floating point processing)


A Graphics Processing Unit is not a CPU!



And there is no GPU without a CPU...

Graphics Processing Unit

- ✓ Originally dedicated to specific operations required by video cards for accelerating graphics, GPU have become flexible **general-purpose** computational engines (GPGPU):
 - ✓ Optimized for **high memory throughput**
 - ✓ Very suitable to **data-parallel processing**
- ✓ Three vendors: Chipzilla (*a.k.a.* Intel), ADI (*i.e.*, AMD), NVIDIA.
- ✓ Traditionally GPU programming has been tricky but **CUDA** and **OpenCL** made it affordable.



Parallel Computer Architectures

1. Memory organization
 - Memory hierarchy
 - Shared versus distributed memory
2. Processor/node architecture
 - Flynn's taxonomy
3. Interconnection network
 - Dynamic versus static networks

07/09/2021

IN4049TU

23


3. Interconnection Networks

Besides the **node architecture** and **memory organization**, **interconnection network** is another important component which characterizes a parallel computer.

Types of Networks:

1. Dynamic
2. Static


Cross-bar networks (dynamic)



September 7, 2021

25


Multi-stage networks (dynamic)



September 7, 2021

26


Perfect Shuffle



September 7, 2021

27

Omega Network (dynamic)




September 7, 2021

28

Complexity Omega Network


- Processors: p
- Stages: $\log p$
- Switches per Stage: $p/2$
- Total number of switches: $\frac{p}{2} \log p$



September 7, 2021

29

Routing in Omega Network



Switching rules:




$\beta_k = 0$: the message is forwarded over the upper link of the switch;
 $\beta_k = 1$: the message is forwarded over the lower link of the switch.

September 7, 2021

30

Collision in Omega Network



Exactly one path for every pair: $s = \alpha_1\alpha_2\dots\alpha_k$ to $t = \beta_1\beta_2\dots\beta_k$

In total $(n/2)\log(n)$ switches $\rightarrow 2^{(n/2)\log(n)} = n^{n/2}$ different switchings compared to $n!$ permutations (for n input to n output), only $n^{n/2}$ of the $n!$ possible permutations can be performed without conflict.

September 7, 2021

31

Dynamic networks (switching network)

- ❖ Flexible in realizing communication of diff. pairs of processors/memories
- ❖ Simple bus type network cheap but not scalable (in performance) to large number of processors
- ❖ Crossbars are very fast but too expensive to scale to large systems
- ❖ Trade-off between cost and performance

32

Static Networks

- Topologies
- Cost models
- Emphasis on nodal-connectivity


07/09/2021

IN4049TU

33

Historical Perspective

- Early machines were:
 - Collection of microprocessors
 - bi-directional queues between neighbors
- Messages were forwarded by processors on path
- Strong emphasis on topology in algorithms



07/09/2021

IN4049TU

34

Network Analogy

- To have a large number of transfers occurring at once, you need a large number of distinct wires
- Networks are like streets
 - link = street
 - switch = intersection
 - distances (hops) = number of blocks traveled
 - routing algorithm = travel plans
- Properties
 - latency: how long to get somewhere in the network
 - bandwidth: how much data can be moved per unit time
 - » limited by the number of wires
 - » and the rate at which each wire can accept data

07/09/2021

IN4049TU

35

Components of a Network

Networks are characterized by

- **Topology** - how things are connected
 - two types of nodes: hosts and switches
- **Routing algorithm** - paths used
 - e.g., all east-west then all north-south (avoids deadlock)
- **Switching strategy**
 - circuit switching: full path reserved for entire message
 - » like the telephone
 - packet switching: message broken into separately-routed packets
 - » like the post office
- **Flow control** - what if there is congestion
 - if two or more messages attempt to use the same channel
 - may stall, move to buffers, reroute, discard, etc.


07/09/2021

IN4049TU

36

Properties of a Network

- **Diameter** is the maximal length of shortest paths between any two nodes in the graph. (another metric: average distance)
- The **bandwidth** of a link is: $w * 1/t$
 - w is the number of wires
 - t is the time per bit
- Effective bandwidth lower due to packet overhead



- **Bisection bandwidth**
 - sum of the minimum number of channels which, if removed, will separate the network into two equal parts
(A network is partitioned if some nodes cannot reach others.)


07/09/2021

IN4049TU

37

Meshes and Tori

- Diameter: $2\sqrt{n}$ (in 2D)
- Bisection bandwidth: \sqrt{n}



- Often used as network in machines
- Generalizes to higher dimensions (Cray T3D used 3D Torus)
- Natural for algorithms with 2D, 3D arrays


07/09/2021

IN4049TU


38

Hypercubes

- Number of nodes $n = 2^d$ for dimension d
 - Diameter: d
 - Bisection bandwidth is $n/2$



- Popular in early machines (Intel iPSC, NCUBE)
 - Lots of clever algorithms
- Greycode addressing
 - each node connected to d others with 1 bit different






07/09/2021

IN4049TU

39

Trees

- Diameter: $\log(n)$
- Bisection bandwidth: 1
- Easy layout as planar graph
- Many tree algorithms (summation)
- Fat trees avoid bisection bandwidth problem
 - more (or wider) links near top
 - example, Thinking Machines CM-5



07/09/2021


IN4049TU

40

Scalable, High Perf. Interconnection Network

- At the core of parallel computer architecture
- Requirements and trade-offs at many levels
 - Elegant mathematical structure
 - Deep relationships to algorithm structure
 - Managing many traffic flows
 - Electrical / Optical link properties
- Little consensus
 - interactions across levels
 - Performance metrics?
 - Cost metrics?
 - Workload?

=> need a holistic understanding



07/09/2021

IN4049TU

41

Example specification: Summit supercomputer

Processors:

4,356 nodes, each with two 22-core Power9 CPUs, and six NVIDIA Tesla V100 GPUs. Total 2,414,592 cores

A Fat-tree network topology

InfiniBand uses a switched fabric topology, as opposed to early shared medium Ethernet.

Implemented using Mellanox 100-Gb/s EDR InfiniBand ConnectX-5 adapters and Switch-IB2 switches

Messages

InfiniBand transmits data in packets of up to 4 KB that are taken together to form a message.

07/09/2021

IN4049TU

42

Example 2 specification: Fukagu Supercomputer

- Configuration: 158,976 Fujitsu A64FX 48C 2.2GHz; total 7299072 cores
- CPU (node) features: 48 cores, 2 SIMD operations/core; memory 32 GB/node with mem. bandwidth 1TB/s.
- Interconnection Network: a 6-dimensional mesh/torus network TOFU, each node with 10 inter-node connections (logically a 3-D torus network)


07/09/2021

IN4049TU


43

TOFU interconnect: 6-D mesh/torus

- ❖ Six coordinate axes: X, Y, Z, A, B, C
 - X, Y, Z: size varies according to system configuration
 - A, B, C: fixed 2×3×2 (building block)
- ❖ TOFU: Torus fusion $X \times Y \times Z \times (2 \times 3 \times 2)$
 - $X \times Y \times Z$ forms a 3-D torus of building blocks.



07/09/2021



TOFU node design: Each pair of adjacent ABC mesh/torus is interconnected via twelve links.

Quiz 1

Submission deadline: 14 September, 15:45 PM

Questions on brightspace in Assignments/Quiz1

Upload your answers to Assignments/Quiz1.