

Multigrid Operators (2)

- **The restriction operator R_i maps $P^{(i)}$ to $P^{(i-1)}$**
 - Restricts problem on fine grid $P^{(i)}$ to coarse grid $P^{(i-1)}$ by sampling or averaging
 - $b_{i-1} = R_i(b_i)$
- **The interpolation operator ln_{i-1} maps an approximate solution x_{i-1} to an x_i**
 - Interpolates solution on coarse grid $P^{(i-1)}$ to fine grid $P^{(i)}$
 - $x_i = ln_{i-1}(x_{i-1})$
- **The solution operator S_i takes $P^{(i)}$ and computes an improved solution $x(i)$ on same grid**
 - Uses “weighted” Jacobi or SOR
 - $x_{i, improved} = S_i(b_i, x_i)$
- **Details of these operators follow after describing overall algorithm**

Multigrid V-Cycle Algorithm (recursive)

Function MGV (b_i , x_i)

... Solve $T_i x_i = b_i$ given b_i and an initial guess for x_i

... return an improved x_i

if ($i = 1$)

compute exact solution x_1 of $P^{(1)}$

return x_1

else

$x_i = S_i(b_i, x_i)$

$r_i = T_i x_i - b_i$

$d_i = \text{In}_{i-1}(\text{MGV} (R_i(r_i) , 0))$


$x_i = x_i - d_i$

$x_i = S_i(b_i, x_i)$

return x_i

only 1 unknown.

Multigrid V-cycle



improve solution by
damping high frequency error,
compute residual,
solve $T_i d_i = r_i$ recursively,
correct fine grid solution,
improve solution again.

Complexity of a V-Cycle

◦ On a serial machine

- Work at each “dot” in the V-cycle is $O(\text{the number of unknowns})$
- Cost of level i is $(2^i-1)^2 = O(4^i)$ (for a 2D grid)
- If finest grid level is m , total time is:

$$\sum_{i=1}^m O(4^i) = O(4^m) = O(\# \text{ unknowns})$$

◦ On a parallel machine (PRAM)

- with one processor per grid point and free communication, each step in the V-cycle takes constant time, $O(1)$
- Total V-cycle time is $O(m) = O(\log \# \text{unknowns})$

Full Multigrid (FMG)

◦ Intuition:

- improve solution by doing multiple V-cycles
- avoid expensive fine-grid (high frequency) cycles
- analysis of why this works is beyond the scope of this class

Function FMG (b_m, x_m)

... return improved x_m given initial guess

compute the exact solution x_1 of $P^{(1)}$


for $i=2$ to m

$$x_i = \text{MGV} (b_i, \text{In}_{i-1} (x_{i-1}))$$

◦ In words:

- Solve the problem with 1 unknown
- Given a solution to the coarser problem, $P^{(i-1)}$, map it to starting guess for $P^{(i)}$
- Solve the finer problem using the Multigrid V-cycle

Full Multigrid Cost Analysis




- **One V for each call to FMG**
 - one also use “W”s and other compositions
- **Serial time:** $\sum_{i=1}^m O(4^i) = O(4^m) = O(\# \text{ unknowns})$
- **PRAM time:** $\sum_{i=1}^m O(i) = O(m^2) = O(\log^2 \# \text{ unknowns})$

Complexity of Solving Poisson's Equation

- **Theorem:** error ε after one FMG call is $\leq c$ times the error before, where $c < 1/2$, and independent of # unknowns
- $\mathbf{x}^k = \text{FMG}(\mathbf{b}, \mathbf{x}^{k-1}) \implies \varepsilon(\mathbf{x}^k) < 1/2 \varepsilon(\mathbf{x}^{k-1})$ (i.e., at least 1 bit per FMG iteration. \mathbf{x}^k = solution after k^{th} FMG iteration)
- **Corollary:** We can make the error $\varepsilon < \text{tol}$, for any fixed tolerance in a fixed number of steps, independent of size of the finest grid
- This is the most important convergence property of MG, distinguishing it from other methods, which converge more slowly for large grids
- Total complexity is just proportional to the cost of one FMG call

The Solution Operator S_i - Details

- The solution operator S_i , is a weighted Jacobi op.
- Consider the 1D problem



- At level i , pure Jacobi replaces:

$$x(j) := 1/2 (x(j-1) + x(j+1) + b(j))$$

- Weighted Jacobi uses:


$$x(j) := 1/3 (x(j-1) + x(j) + x(j+1) + b(j))$$

- In 2D, similar average of nearest neighbors

The Restriction Operator R_i - Details



- The restriction operator R_i takes
 - a problem $P^{(i)}$ with RHS b_i and
 - maps it to a coarser problem $P^{(i-1)}$ with RHS b_{i-1}
- In 1D, average values of neighbors
 - $x_{\text{coarse}}(j) = 1/4 * x_{\text{fine}}(j-1) + 1/2 * x_{\text{fine}}(j) + 1/4 * x_{\text{fine}}(j+1)$

The Restriction Operator R_i - Details




° In 2D, average with all 8 neighbors (N,S,E,W,NE,NW,SE,SW)

Interpolation Operator In_i



- ° In 2D, interpolation requires averaging with 4 nearest neighbors (NW,SW,NE,SE)



Performance Model of parallel 2D Multigrid

- Assume 2^m+1 by 2^m+1 grid of unknowns, $n = 2^m+1$, $N = n^2$
- Assume $p = 4^k$ processors, arranged in 2^k by 2^k grid
 - Each processor starts with 2^{m-k} by 2^{m-k} subgrid of unknowns
- Consider V-cycle starting at level m
 - At levels m through k of V-cycle, each processor does some work
 - At levels $k-1$ through 1 , some processors are idle, because a 2^{k-1} by 2^{k-1} grid of unknowns cannot occupy each processor

Performance Model of parallel 2D Multigrid (2)

◦ Cost of one level (j) or $P^{(j)}$ in V-cycle

- If level $j \geq k$, then cost =

$O(4^{j-k})$ Flops, proportional to the number of grid points/processor
+ $O(1) \alpha$ Send a constant # messages to neighbors
+ $O(2^{j-k}) \beta$ Number of words sent

- If level $j < k$, then cost =

$O(1)$ Flops, proportional to the number of grid points/processor
+ $O(1) \alpha$ Send a constant # messages to neighbors
+ $O(1) \beta$ Number of words sent

◦ Sum over all levels in all V-cycles in FMG to get complexity

Comparison of Methods (using p processors)

	# Flops	# Messages	# Words sent
MG	$N/p + \log p * \log N$	$(\log N)^2$	$(N/p)^{1/2} + \log p * \log N$
FFT	$N \log N / p$	$p^{1/2}$	N/p
SOR	$N^{3/2}/p$	$N^{1/2}$	N/p

- **SOR is slower than others on all counts**
- Flops for MG and FFT depends on accuracy of MG
- **MG communicates less total data (bandwidth)**
- **Total messages (latency) depends ...**
 - This coarse analysis can't say whether MG or FFT is better when $\alpha \gg \beta$

Practicalities

- In practice, we don't go all the way to $P^{(1)}$
- In sequential code, the coarsest grids are negligibly cheap, but on a parallel machine they are not.

- Consider 1000 points per processor
- In 2D, the surface to communicate is $4 \cdot \sqrt{1000} \approx 128$, or 13%
- In 3D, the surface is $1000 - 8^3 \approx 500$, or 50%
- Data locality ratio α (large α is preferred):

$$\alpha = \frac{\text{computation time between two communication steps}}{\text{communication time}}$$

- **Apply communication avoiding Jacobi iterations!**
- **Dealing with coarse meshes efficiently**
 - Should we switch to using fewer processors on coarse meshes?
 - Should we switch to another solver on coarse meshes?