

Performance metrics

A pragmatic classification:

- Users :
 - How fast is my application?
 - *execution time*
- Developers:
 - How close to the absolute best can I be?
 - *estimate absolute best*
 - *compute performance gain (speed-up)*
- Budget-holders:
 - How much of my infrastructure am I using?
 - *efficiency*
 - *utilization*

Performance “actions”

- Performance measurement
 - *measure execution time*
 - *derive metrics such as speed-up, throughput, bandwidth*
 - *platform and application implementation are available*
 - *data-sets are available*
- Performance analysis
 - estimate performance bounds
 - *performance bounds are typically worst-case, best-case, average-case scenarios*
 - platform and application are available/models
 - data - sets are available/models
- Performance prediction
 - estimate application behavior
 - platform and application are models
 - data - sets are real

Performance Metrics

- Serial execution time: T_S
- Parallel execution time: T_P
- Overhead (p is # compute units) $T_O = p \cdot T_P - T_S$
 - *ideal case: $T_O = 0$ (perfect linear speed-up)*

Total overhead relevant for
dedicated parallel processing

- Speedup $S = \frac{T_{serial_best}}{T_P}$

Overhead may
depend on p !

Relative versus true speedup: using $T_P(P=1)$ instead of T_{serial_best}

Superlinear speed-up is sometimes possible:
cache effects / memory sizes


Efficiency

$$E = \frac{S}{p}$$

$$E = \frac{S}{p} = \frac{T_S}{p \cdot T_P}$$

$$T_O = p \cdot T_P - T_S$$

$$E = \frac{1}{1 + \frac{T_O}{T_S}}$$




Sources of overhead in parallel programs

- Inter-process interaction
 - *communication => idling*
 - *synchronization => serialization*
- Load imbalance
 - *un-even workloads => idling*
- Additional computations, such as
 - *memory allocation*
 - *data partitioning*
 - *managing the parallelism*
 - ...

Recap: Amdahl's law – fixed problem size

- Every application has an intrinsically sequential part
- Amdahl's law:
 - *let s be the fraction of work that is sequential, then $(1-s)$ is the fraction that is parallelizable*
 - *p = number of processors*
 - *S = Speedup*



$$\begin{aligned} S &= T_{seq}/T_{par} \\ &= 1/(s + (1-s)/p) \\ &\leq 1/s \end{aligned}$$

Speedup is bounded by the sequential fraction.

Isoefficiency function

What is a 'good' efficiency of a parallel program? → No simple answer

Emphasis on scalability of a parallel algorithm: can a program retain its efficiency when #processors and problem size increase?

$$T_o = p \cdot T_p - W \quad (\text{definition of overhead})$$


and

$$S = \frac{W}{T_p} \quad \Rightarrow \quad S = \frac{W \cdot p}{W + T_o} \quad \Rightarrow \quad E = \frac{1}{1 + T_o/W}$$


Conclusions:

1. $E=1$ when there is no overhead
2. E is fixed iff T_o/W is fixed

Example: stencil-type computations (1/3)



$(n+2) \times (n+2)$ grid
given boundary values



column-wise
data distribution

```
a[i,j]=0.25*(a[i,j-1]+a[i-1,j]+a[i,j+1]+a[i+1,j]);
```

$$T_S = 4t_{op} \cdot n^2$$

$$T_{comm} = 2n \cdot t_{data}$$

$$T_{calc} = 4t_{op} \cdot (n \cdot n/p) = 4t_{op}n^2/p$$

$$T_P = 4t_{op} \cdot n^2/p + 2n \cdot t_{data}$$

Example: stencil-type computations (2/3)

$$S = \frac{t_{op}n^2 \cdot p}{t_{op}n^2 + p \cdot n \cdot t_{data}/2}$$

$$E = \frac{S}{p} = \frac{1}{1 + (p/n)(t_{data}/2t_{op})}$$

- So p must be small relative to n for efficiency
- Efficiency stays constant as long as p/n is constant

Hardware Performance metrics

- Clock frequency [GHz] = absolute hardware speed
 - *memories, CPUs, interconnects*
- Operational speed [GFLOPs]
 - *how many operations per cycle a machine can do*
- Memory bandwidth (BW) [GB/s]
 - *differs a lot between different memories on chip*
 - *remember? Slow memory is large, fast memory is small ...*
- Power [Watt]
- Derived metrics
 - *normalized for comparison purposes ...*
 - *FLOPs/Byte, FLOPs/Watt, ...*

Theoretical peak performance

$$\text{Peak} = \text{chips} * \text{cores} * \text{vectorWidth} * \\ \text{FLOPs/cycle} * \text{clockFrequency}$$

- *cores = real cores, hardware threads, or ALUs, depending on the architecture*

Examples from DAS-4:

- Intel Core i7 CPU

$$2 \text{ chips} * 4 \text{ cores} * 4\text{-way vectors} * 2 \text{ FLOPs/cycle} * 2.4 \text{ GHz} = 154 \text{ GFLOPs}$$

- NVIDIA GTX 580 GPU

$$1 \text{ chip} * 16 \text{ SMs} * 32 \text{ cores} * 2 \text{ FLOPs/cycle} * 1.544 \text{ GHz} = 1581 \text{ GFLOPs}$$

- ATI AMD Radeon HD 6970 GPU

$$1 \text{ chip} * 24 \text{ SIMD engines} * 16 \text{ cores} * 4\text{-way vectors} * 2 \text{ FLOPs/cycle} \\ * 0.880 \text{ GHz} = 2703 \text{ GFLOPs}$$

DRAM Memory bandwidth (off-chip)

Throughput = memory bus frequency * bits per cycle * bus width

- *memory clock is not the CPU clock (typically lower)*
- *divide by 8 to get B/s*

Examples:

- Intel Core i7 DDR3: $1.333 \text{ GHz} * 2 * 64 = 21 \text{ GB/s}$
- NVIDIA GTX 580 GDDR5: $1.002 \text{ GHz} * 4 * 384 = 192 \text{ GB/s}$
- ATI HD 6970 GDDR5: $1.375 \text{ GHz} * 4 * 256 = 176 \text{ GB/s}$

Power

- Chip manufacturers specify Thermal Design Power (TDP)
 - *some definition of maximum power consumption ...*
- We can measure dissipated power
 - *whole system*
 - *typically (much) lower than TDP*
- Power efficiency: FLOPs / Watt
- Examples (with theoretical peak and TDP)
 - Intel Core i7: $154 / 160 = 1.0 \text{ GFLOPs/W}$
 - NVIDIA GTX 580: $1581 / 244 = 6.3 \text{ GFLOPs/W}$
 - ATI HD 6970: $2703 / 250 = 10.8 \text{ GFLOPs/W}$

Software metrics (3 P's)

Performance metrics

- Execution time
 - *Derive speed-up vs. best available sequential performance*
- Achieved GFLOPs:
 - *Count (FL)OPs, divide by execution time => FLOPS/s*
 - *Derive computational efficiency (i.e., utilization) = $\frac{\text{Achieved FLOPs}}{\text{Peak FLOPs}}$*
- Achieved GB/s:
 - *Count memory OPs, divide by execution time => B/s*
 - *Derive memory efficiency (i.e., utilization) = $\frac{\text{Achieved GB/s}}{\text{Peak GB/s}}$*

Productivity and Portability metrics

- Programmability
- Production costs
- Maintenance costs

Attainable performance

- Attainable GFlops/sec
= \min (Peak Floating -Point Performance,
Peak Memory Bandwidth * Operational Intensity)
- To translate:
 - if an application is compute-bound =>
performance is limited by peak performance
 - if an application is memory-bound =>
performance is limited by the load it puts on the memory system

Use the Roofline model

Determine what to do first to gain performance

- increase memory streaming rate (fights mem-boundness)
 - *GPU: memory coalescing*
 - *CPUs: better caching*
- apply in-core optimizations (fights compute-boundness)
 - *vectorization*
- increase arithmetic intensity (fights mem-boundness)
 - *change your algorithm*
 - *think of new ways to reuse the data*