

# Particle Simulation, bottlenecks

---

## ◦ Computational bottlenecks?

- force calculation; contains a nested loop **Yes**
- move operation; fixed amount of work per particle **No**
- calculation of properties **Possible**

## ◦ For high performance, focus on bottleneck force calculation



## ◦ Types of force to distinguish (in practice)

**(in parallel code)**

- External force **Trivial**
  - Is independent on the other particles
- Short-range force **Easy**
  - Depends on particles within fixed range
- Long-range force **Difficult**
  - Depends on all other particles

# Example of an Adaptive Quad Tree

Adaptive quadtree where no square contains more than 1 particle



Child nodes enumerated counterclockwise from SW corner, empty ones are excluded

# Adaptive Quad Tree Building Cost

---

° **Cost**     $\leq N * \text{maximum cost of a Quad\_Tree\_Insert}$   
               $= O(N * \text{maximum depth of QuadTree})$

° **Uniform distribution:**

**depth of QuadTree** =  $O(\log N)$ ,  
      **so Cost** =  $O(N \log N)$

° **Arbitrary distribution:**

**depth of Quad Tree** =  $O(b) = O(\# \text{ bits in particle coordinates})$ ,  
      **so Cost** =  $O(b N)$

– **Note:** depth of QuadTree  $b \geq 1 + 4 \log N$

# Barnes-Hut Algorithm

---

## ◦ High Level description:

- 1) **Build the QuadTree using QuadTreeBuild**  
... already described, cost =  $O(N \log N)$  or  $O(b N)$
- 2) **For each node, a sub-square in the QuadTree,**  
**Compute the CM and total mass (TM) of all the particles it contains**  
... “post order traversal” of QuadTree, cost =  $O(N \log N)$  or  $O(b N)$
- 3) **For each particle, traverse the QuadTree to compute the force on it,**  
**using the CM and TM of “distant” sub-squares**  
... core of algorithm  
... cost depends on accuracy desired but still  $O(N \log N)$  or  $O(bN)$


**Total cost: still  $O(N \log N)$  or  $O(bN)$**

→ Achieved order reduction per time step:  $O(N^2)$  to  $O(N \log N)$

## Step 3 of BH: Analysis

- **Correctness: recursive accumulation of force from each subtree**
  - Each particle is accounted for exactly once, whether it is in a leaf or other node
- **Complexity analysis**
  - Cost of  $\text{TreeForce}(k, \text{root}) = O(\text{depth in QuadTree of leaf containing } k)$
  - “Proof”; Example: Assume  $\theta = 1$ 
    - For each undivided node, see fig., (except one containing  $k$ ),  $D/r < 1 < \theta$
    - There are only 3 nodes to consider at each level of the QuadTree, see fig.
    - There is  $O(1)$  work per node
    - Cost =  $O(\text{level of } k)$
  - Total cost =  $O(\sum_k \text{level of } k) = O(N \log N)$ 
    - Strongly depends on  $\theta$

Sample Barnes-Hut Force calculation  
For particle in lower right corner  
Assuming  $\theta > 1$



# Fast Multiple Method (FMM)

---

- “A fast algorithm for particle simulation”, L. Greengard and V. Rokhlin, J. Comp. Phys. V. 73, 1987, many later papers
  
- **Differences from Barnes-Hut**
  - FMM computes the *potential* at every point, not just the force
  - FMM uses more information in each box than the CM and TM, so it is both more accurate and more expensive
  
  - In compensation, FMM accesses a fixed set of boxes at every level, independent of  $D/r$
  
  - BH uses fixed information (CM and TM) in every box, but # boxes increases with accuracy.
  - FMM uses a fixed # boxes, but the amount of information per box increase with accuracy.

# Parallelizing Hierarchical N-Body codes

---

- **Barnes-Hut, FMM and related algorithms have similar computational structure:**
  - 1) **Build the QuadTree**
  - 2) **Traverse QuadTree from leaves to root and build outer expansions (just (TM,CM) for Barnes-Hut)**
  - 3) **Traverse QuadTree from root to leaves and build any inner expansions**
  - 4) **Traverse QuadTree to accumulate forces for each particle**
- **QuadTree changes dynamically when the particles move, so the tree has to be rebuilt (or adjusted) every time step.**
- **But: No doubly nested loop over all particles anywhere in the algorithm**
- **All 4 phases have to be parallelized efficiently.**

# Parallelizing Hierarchical N-Body codes

---

## ◦ General idea: Domain decomposition

- Assign regions of space to each processor
- Regions may have different size or shape, to get a good load balance
  - Each region will have about  $N/p$  particles
- Each processor will store part of QuadTree containing all particles (=leaves) in its region, and their ancestors in QuadTree
  - Root of tree and some generations stored by all processors, nodes may also be shared
- Each processor will also store adjoining parts of QuadTree needed to compute forces for particles it owns
  - Subset of QuadTree needed by a processor called the **Locally Essential Tree (LET)**
- Given the LET, all force accumulations (step 4)) are done in parallel, without communication

## ◦ Coarse grained parallelism

- Each domain solves its own N-body problem, but somehow has to take into account the effect of all the other particles as well; like the ghost-points in the Poisson problem. Here those particles generate some “background” potential (in FMM)




# Load Balancing Scheme 1

---

- **Orthogonal Recursive Bisection (ORB) of space**
  - Warren and Salmon, Supercomputing 92
- **Recursively split region along axes into regions containing equal numbers of particles**
  - Particles are grouped in rectangular regions; may be very elongated
  - No relation with tree

**Partitioning for 16 procs:**

Orthogonal Recursive Bisection



## Load Balancing Scheme 2



---

- **Idea: Partition QuadTree instead of space**
  - Estimate work for each node, call total work  $W$
  - Arrange nodes of QuadTree in **some** linear order (lots of choices)
  - Assign contiguous groups of nodes with work  $W/p$  to processors
- **Method called: Costzones or Hashed Tree**
  - J.P. Singh, PhD thesis, Stanford, 1993
  - Warren and Salmon, Supercomputing 93

## Load Balancing Scheme 2

- Make sure that neighboring leaves in the tree are also neighboring in space
  - Which of the 4 children of a node is “first” depends on the position of the node
  - Orientation changes: clockwise / counter-clockwise

Using costzones to layout a quadtree on 4 processors  
Leaves are color coded by processor color



# Implementing Costzones

---

## ◦ Random Sampling

- All processors own some particles,
- All processors send a small random sample of their particles to Processor 1
- Processor 1
  - builds small Quadtree serially,
  - determines its Costzones, and
  - broadcasts them to all processors
- Other processors build the part of Quadtree they are assigned by these Costzones

## ◦ All processors know all Costzones

## ◦ This is needed later to compute LET's

# Locally Essential Trees (LETs)

---

- **Warren and Salmon, 1992; Liu and Bhatt, 1994**
- **Definition:**
  - **A LET of a process is that part of the Tree that is necessary to compute the force on the particles that are owned by that process.**
- **Information about nodes near the root of the tree is present in all processes**
- **Information about nodes near some leaves of the tree that are all owned by a single process is needed in one or a few processes**

# **Computing Locally Essential Trees (LETs)**

---

- **Warren and Salmon, 1992; Liu and Bhatt, 1994**
- **Every processor needs a subset of the whole QuadTree, called the LET, to compute the force on all particles it owns**
- **Shared Memory**
  - Receiver driven protocol
  - Each processor reads part of QuadTree it needs from shared memory on demand, keeps it in cache
  - Drawback: cache memory appears to need to grow proportionally to  $P$  to remain scalable
- **Distributed Memory**
  - Sender driven protocol
  - Each processor decides which other processors need parts of its local subset of the Quadtree, and sends these subsets

# Locally Essential Trees in Distributed Memory

---

## ◦ Barnes-Hut

### Which nodes are needed?

- Let  $j$  and  $k$  be processes,  $n$  a node on process  $j$
- Let  $D(n)$  be the side length of  $n$
- Let  $r_k(n)$  be the shortest distance from  $n$  to any point owned by  $k$
- If either
  - (1)  $D(n)/r_k(n) < \theta$  and  $D(\text{parent}(n))/r_k(\text{parent}(n)) \geq \theta$ , or
  - (2)  $D(n)/r_k(n) \geq \theta$then node  $n$  is part of  $k$ 's LET, and so process  $j$  should send  $n$  to  $k$
- Condition (1) means (TM,CM) of  $n$  can be used on process  $k$ , but this is not true of any ancestor
- Condition (2) means that we need the ancestors of type (1) nodes too