
Performance Analysis of Parallel Programs

Lecture 9

Introduction to High Performance Computing

IN4049 TUDelft

2021 / 2022

About Performance

- Performance metrics
- Performance analysis, estimation, and prediction
 - theoretical performance analysis
 - *speed-up*
 - *efficiency*
 - practical performance analysis
 - *hardware*
 - *software*

Performance metrics

A pragmatic classification:

- Users :
 - How fast is my application?
 - *execution time*
- Developers:
 - How close to the absolute best can I be?
 - *estimate absolute best*
 - *compute performance gain (speed-up)*
- Budget-holders:
 - How much of my infrastructure am I using?
 - *efficiency*
 - *utilization*

Performance “actions”

- Performance measurement
 - *measure execution time*
 - *derive metrics such as speed-up, throughput, bandwidth*
 - *platform and application implementation are available*
 - *data-sets are available*
- Performance analysis
 - estimate performance bounds
 - *performance bounds are typically worst-case, best-case, average-case scenarios*
 - platform and application are available/models
 - data-sets are available/models
- Performance prediction
 - estimate application behavior
 - platform and application are models
 - data-sets are real

Theoretical performance analysis

Performance Metrics

- Serial execution time: T_S
 - Parallel execution time: T_P
 - **Overhead** (p is # compute units) $T_O = p \cdot T_P - T_S$
 - *ideal case: $T_O = 0$ (perfect linear speed-up)*
- Total overhead relevant for dedicated parallel processing

- **Speedup** $S = \frac{T_{serial_best}}{T_P}$

Overhead may
depend on p !

Relative versus true speedup: using $T_P(P=1)$ instead of T_{serial_best}

Superlinear speed-up is sometimes possible:
cache effects / memory sizes

Sub-/Super-/linear Speedup

$$S_{linear} = p$$

Linear Speedup

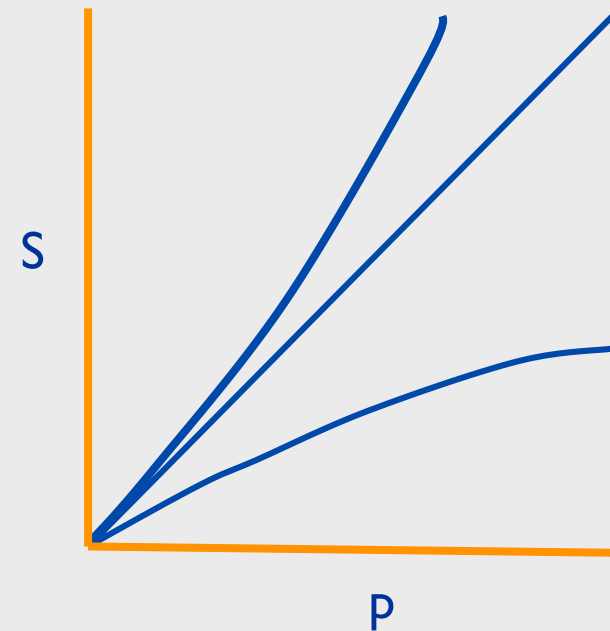
$$S_{sublinear} < p$$

Sublinear speedup

$$S_{superlinear} > p$$

Superlinear speedup

E.g., a small portion of data can yield a higher cache hit rate



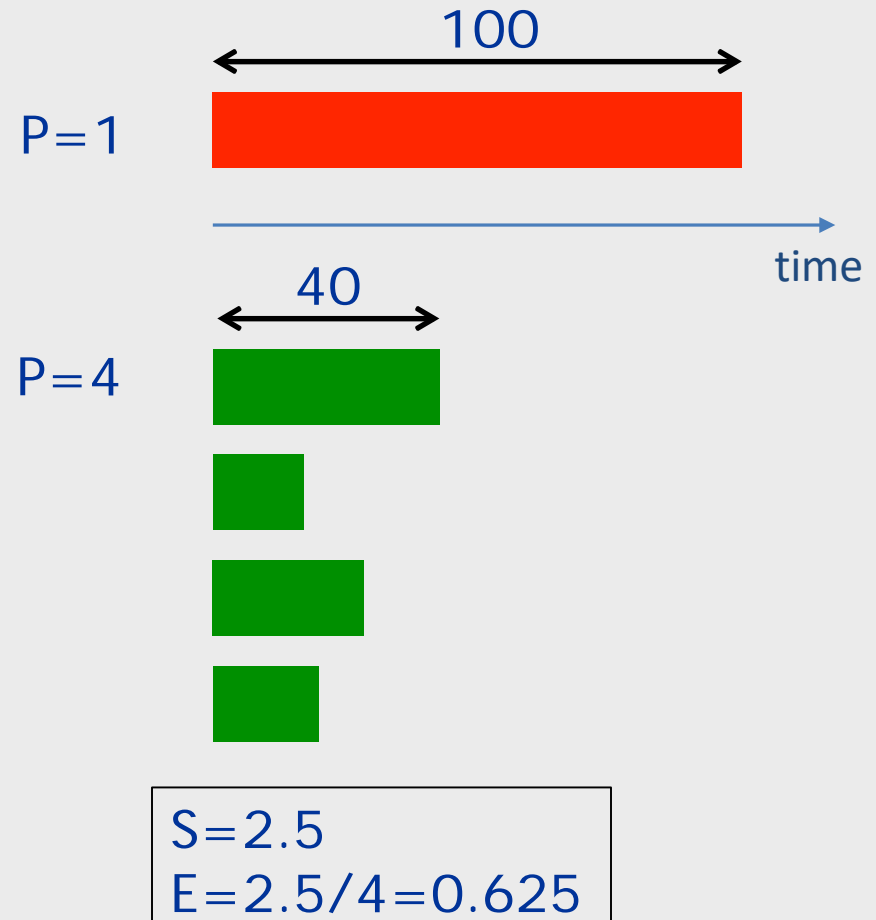
Efficiency

$$E = \frac{S}{p}$$

$$E = \frac{S}{p} = \frac{T_s}{p \cdot T_p}$$

$$T_O = p \cdot T_P - T_S$$

$$E = \frac{1}{1 + \frac{T_o}{T_s}}$$



Sources of overhead in parallel programs

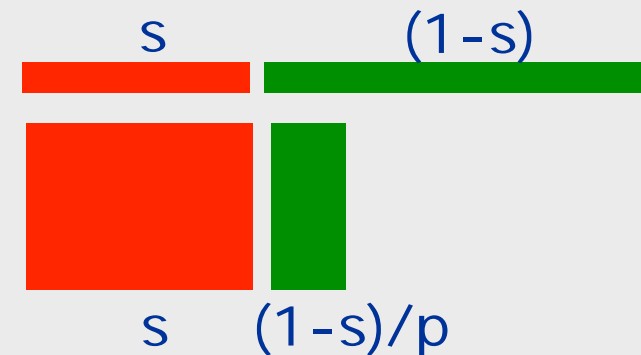
- Inter-process interaction
 - *communication* => *idling*
 - *synchronization* => *serialization*
- Load imbalance
 - *un-even workloads* => *idling*
- Additional computations, such as
 - *memory allocation*
 - *data partitioning*
 - *managing the parallelism*
 - ...

What is problem size?

- Intuitively: **size of the input** ... but
 - *multiple inputs are a problem*
 - *does not characterize the dependency on the computation*
 - example: MVP: $A(m,n) * y(n)$
- Better definition:
 - *the problem size (W) is equal to **the number of basic computation steps** in the best sequential algorithm*
 - example: MVP: $W = \Theta(m*n)$
- We assume further that **problem size W** = T_s

Recap: Amdahl's law – fixed problem size

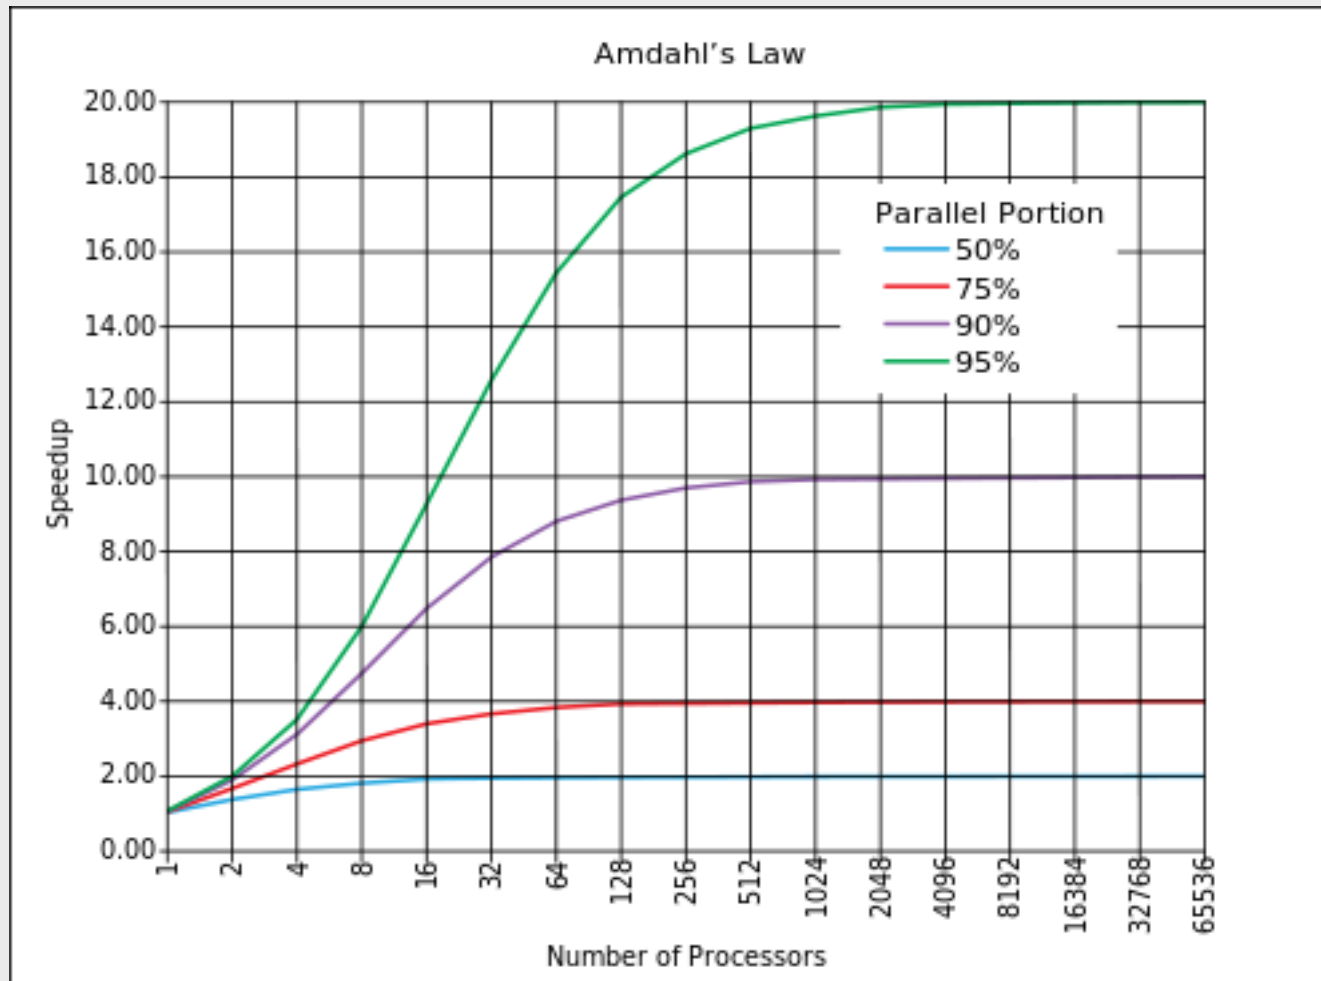
- Every application has an intrinsically **sequential part**
- **Amdahl's law:**
 - let s be the fraction of work that is sequential, then $(1-s)$ is the fraction that is parallelizable
 - p = number of processors
 - S = Speedup



$$\begin{aligned} S &= T_{seq}/T_{par} \\ &= 1/(s + (1-s)/p) \\ &\leq 1/s \end{aligned}$$

Speedup is bounded by the sequential fraction.

Amdahl's Law: $\text{max speedup} < 1/s$



Isoefficiency function

What is a 'good' efficiency of a parallel program? → No simple answer

Emphasis on **scalability** of a parallel algorithm: can a program retain its efficiency when #processors and problem size increase?

$$T_o = p \cdot T_P - W \quad (\text{definition of overhead})$$

and

$$S = \frac{W}{T_P} \quad \Rightarrow \quad S = \frac{W \cdot p}{W + T_o} \quad \Rightarrow \quad E = \frac{1}{1 + T_o/W}$$

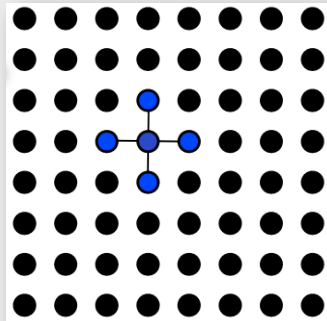
Conclusions:

1. $E=1$ when there is no overhead
2. E is fixed iff T_o/W is fixed

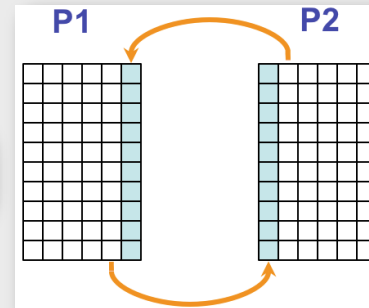
Simple performance modeling

- **Model computation**
 - *count number of operations*
 - *assume flat memory model*
- **Model communication**
 - *typically simple model, linear with the number of data items communicated for large volumes of data*
 - *only model explicit the distant communication*
- Assume:
 - $T_s = \text{number_ops} * t_{op}$
 - $T_p = (\text{number_ops}/p) * t_{op} + T_{comm}$
 - $T_{comm} = \text{number_comm} * t_{comm} + t_{setup}$ (large setup time as will be seen later, we ignore it for simplicity now)

Example: stencil-type computations (1/3)



$(n+2) \times (n+2)$ grid
given boundary values



column-wise
data distribution

```
a[i,j]=0.25*(a[i,j-1]+a[i-1,j]+a[i,j+1]+a[i+1,j]);
```

$$T_S = 4t_{op} \cdot n^2$$

$$T_{comm} = 2n \cdot t_{data}$$

$$T_{calc} = 4t_{op} \cdot (n \cdot n/p) = 4t_{op}n^2/p$$

$$T_P = 4t_{op} \cdot n^2/p + 2n \cdot t_{data}$$

Example: stencil-type computations (2/3)

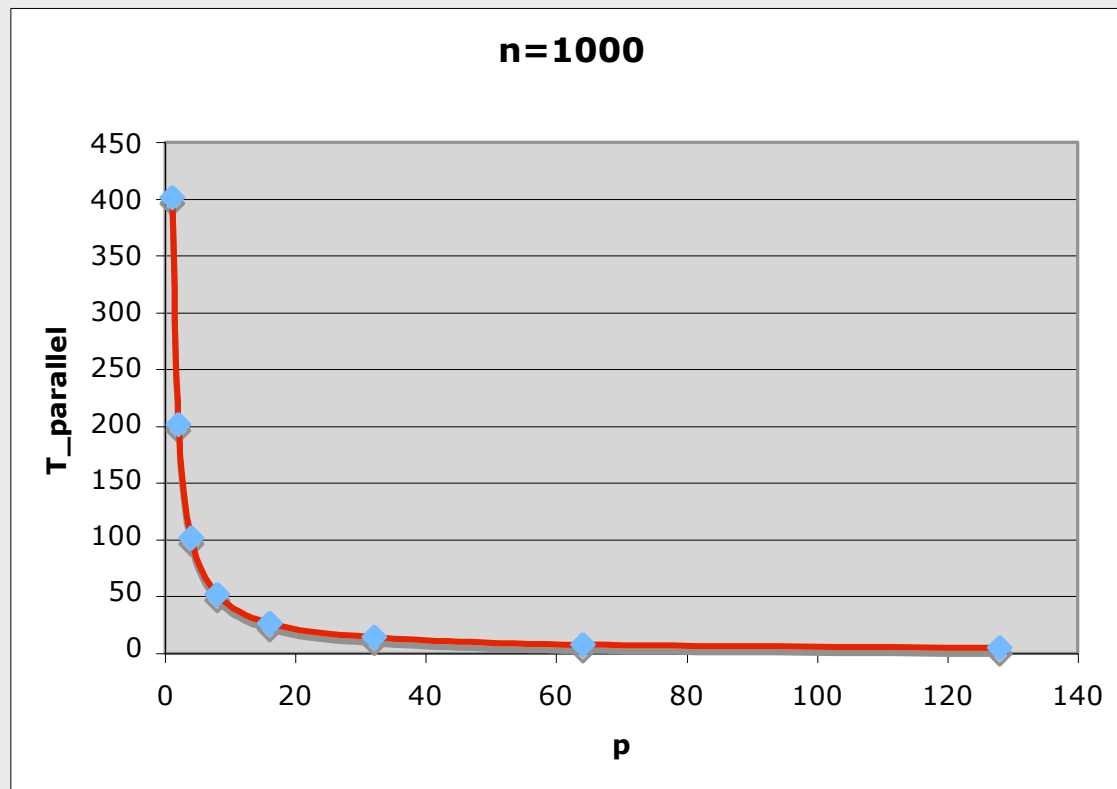
$$S = \frac{t_{op}n^2 \cdot p}{t_{op}n^2 + p \cdot n \cdot t_{data}/2}$$

$$E = \frac{S}{p} = \frac{1}{1 + (p/n)(t_{data}/2t_{op})}$$

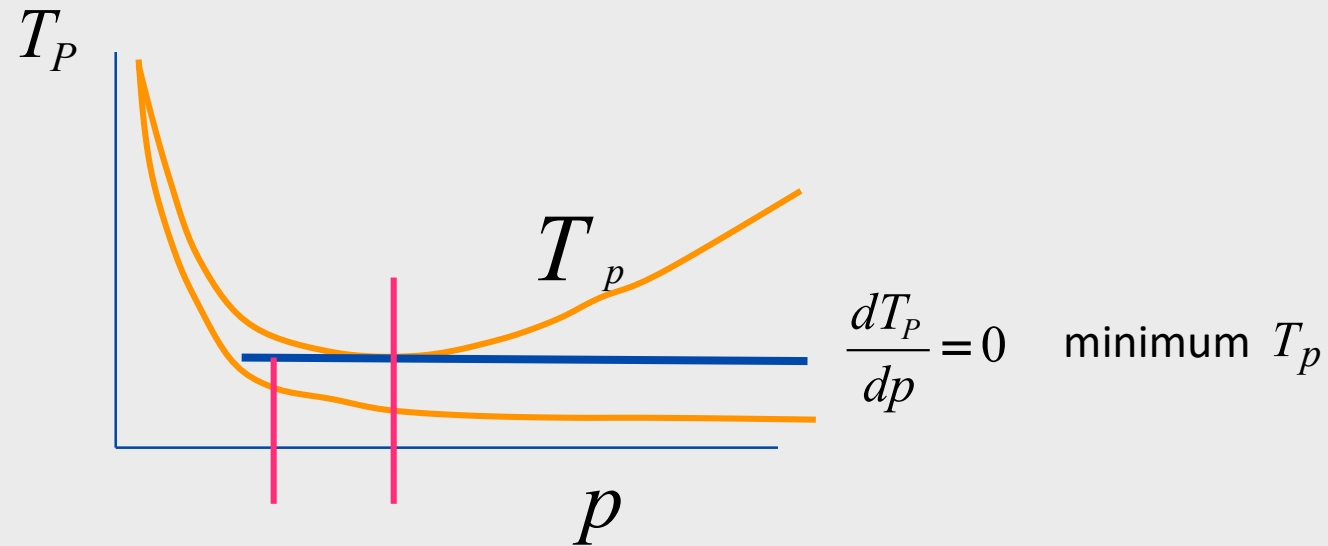
- So p must be small relative to n for efficiency
- Efficiency stays constant as long as p/n is constant

Example: stencil-type computations (3/3)

Suppose $t_{\text{data}} = 20 \cdot t_{\text{op}}$, $n = 1000$



Typical Performance Curve



$$E = 90\%$$

For stencil example
($n=1000$): $p \approx 22$

Performance evaluation: practical examples

Hardware Performance metrics

- **Clock frequency** [GHz] = absolute hardware speed
 - *memories, CPUs, interconnects*
- **Operational speed** [GFLOPs]
 - *how many operations per cycle a machine can do*
- **Memory bandwidth** (BW) [GB/s]
 - *differs a lot between different memories on chip*
 - *remember? Slow memory is large, fast memory is small ...*
- **Power** [Watt]
- **Derived metrics**
 - *normalized for comparison purposes ...*
 - *FLOPs/Byte, FLOPs/Watt, ...*

Theoretical peak performance

Peak = chips * cores * vectorWidth *

FLOPs/cycle * clockFrequency

- *cores = real cores, hardware threads, or ALUs, depending on the architecture*

Examples from DAS-4:

- Intel Core i7 CPU

2 chips * 4 cores * 4-way vectors * 2 FLOPs/cycle * 2.4 GHz = **154 GFLOPs**

- NVIDIA GTX 580 GPU

1 chip * 16 SMs * 32 cores * 2 FLOPs/cycle * 1.544 GHz = **1581 GFLOPs**

- ATI AMD Radeon HD 6970 GPU

1 chip * 24 SIMD engines * 16 cores * 4-way vectors * 2 FLOPs/cycle
* 0.880 GHz = **2703 GFLOPs**

DRAM Memory bandwidth (off-chip)

Throughput = memory bus frequency * bits per cycle * bus width

- *memory clock is not the CPU clock (typically lower)*
- *divide by 8 to get B/s*

Examples:

- Intel Core i7 DDR3: $1.333 \text{ GHz} * 2 * 64 = 21 \text{ GB/s}$
- NVIDIA GTX 580 GDDR5: $1.002 \text{ GHz} * 4 * 384 = 192 \text{ GB/s}$
- ATI HD 6970 GDDR5: $1.375 \text{ GHz} * 4 * 256 = 176 \text{ GB/s}$

Memory bandwidths

On-chip memory can be orders of magnitude faster

- Registers, shared memory, caches, ...
 - *e.g., AMD HD 7970 L1 cache achieves 2 TB/s*

Other memories: depends on the interconnect

- Intel's QPI (Quick Path Interconnect) : **25.6 GB/s**
- AMD's HT3 (Hyper Transport 3) : **19.2 GB/s**
- Accelerators: PCI-e 2.0 : **8.0 GB/s**

Power

- Chip manufacturers specify **Thermal Design Power (TDP)**
 - *some definition of maximum power consumption ...*
- We can measure dissipated power
 - *whole system*
 - *typically (much) lower than TDP*
- **Power efficiency**: FLOPs / Watt
- Examples (with theoretical peak and TDP)
 - Intel Core i7: $154 / 160 =$ **1.0 GFLOPs/W**
 - NVIDIA GTX 580: $1581 / 244 =$ **6.3 GFLOPs/W**
 - ATI HD 6970: $2703 / 250 =$ **10.8 GFLOPs/W**

Summary

	Cores	Threads/ ALUs	GFLOPS	Bandwidth	FLOPs/ Byte
Sun Niagara 2	8	64	11.2	76	0.1
IBM BG/P	4	8	13.6	13.6	1.0
IBM Power 7	8	32	265	68	3.9
Intel Core i7	4	16	85	25.6	3.3
AMD Barcelona	4	8	37	21.4	1.7
AMD Istanbul	6	6	62.4	25.6	2.4
AMD Magny- Cours	12	12	125	25.6	4.9
Cell/B.E.	8	8	205	25.6	8.0
NVIDIA GTX 580	16	512	1581	192	8.2
NVIDIA GTX 680	8	1536	3090	192	16.1
AMD HD 6970	384	1536	2703	176	15.4
AMD HD 7970	32	2048	3789	264	14.4

Absolute hardware performance

Only achieved in the optimal conditions:

- processing units 100% used
- all parallelism 100% exploited
- all data transfers at maximum bandwidth
- but
 - *no application can achieve all of this*
 - *even difficult to write micro-benchmarks*
- hardware catalogue values are not realistic estimates of performance upper bounds

We need realistic estimates of what a platform can do
when given a **realistic workload** => **Roofline** model

Software metrics (3 P's)

Performance metrics

- Execution time
 - *Derive speed-up vs. best available sequential performance*
- Achieved GFLOPs:
 - *Count (FL)OPs, divide by execution time => FLOPS/s*
 - *Derive computational efficiency (i.e., utilization) = $\frac{\text{Achieved FLOPs}}{\text{Peak FLOPs}}$*
- Achieved GB/s:
 - *Count memory OPs, divide by execution time => B/s*
 - *Derive memory efficiency (i.e., utilization) = $\frac{\text{Achieved GB/s}}{\text{Peak GB/s}}$*

Productivity and Portability metrics

- Programmability
- Production costs
- Maintenance costs

Arithmetic intensity

- The (average) number of **arithmetic (floating point) operations** executed by a kernel per byte of memory accessed
- Ignore “overheads”
 - *loop counters*
 - *array index calculations*
 - ...

Convert color into grayscale

```
for (int y = 0; y < height; y++) {  
    for (int x = 0; x < width; x++) {  
        Pixel pixel = RGB[y][x];  
        gray[y][x] =  
            0.30 * pixel.R  
            + 0.59 * pixel.G  
            + 0.11 * pixel.B;  
    }  
}
```

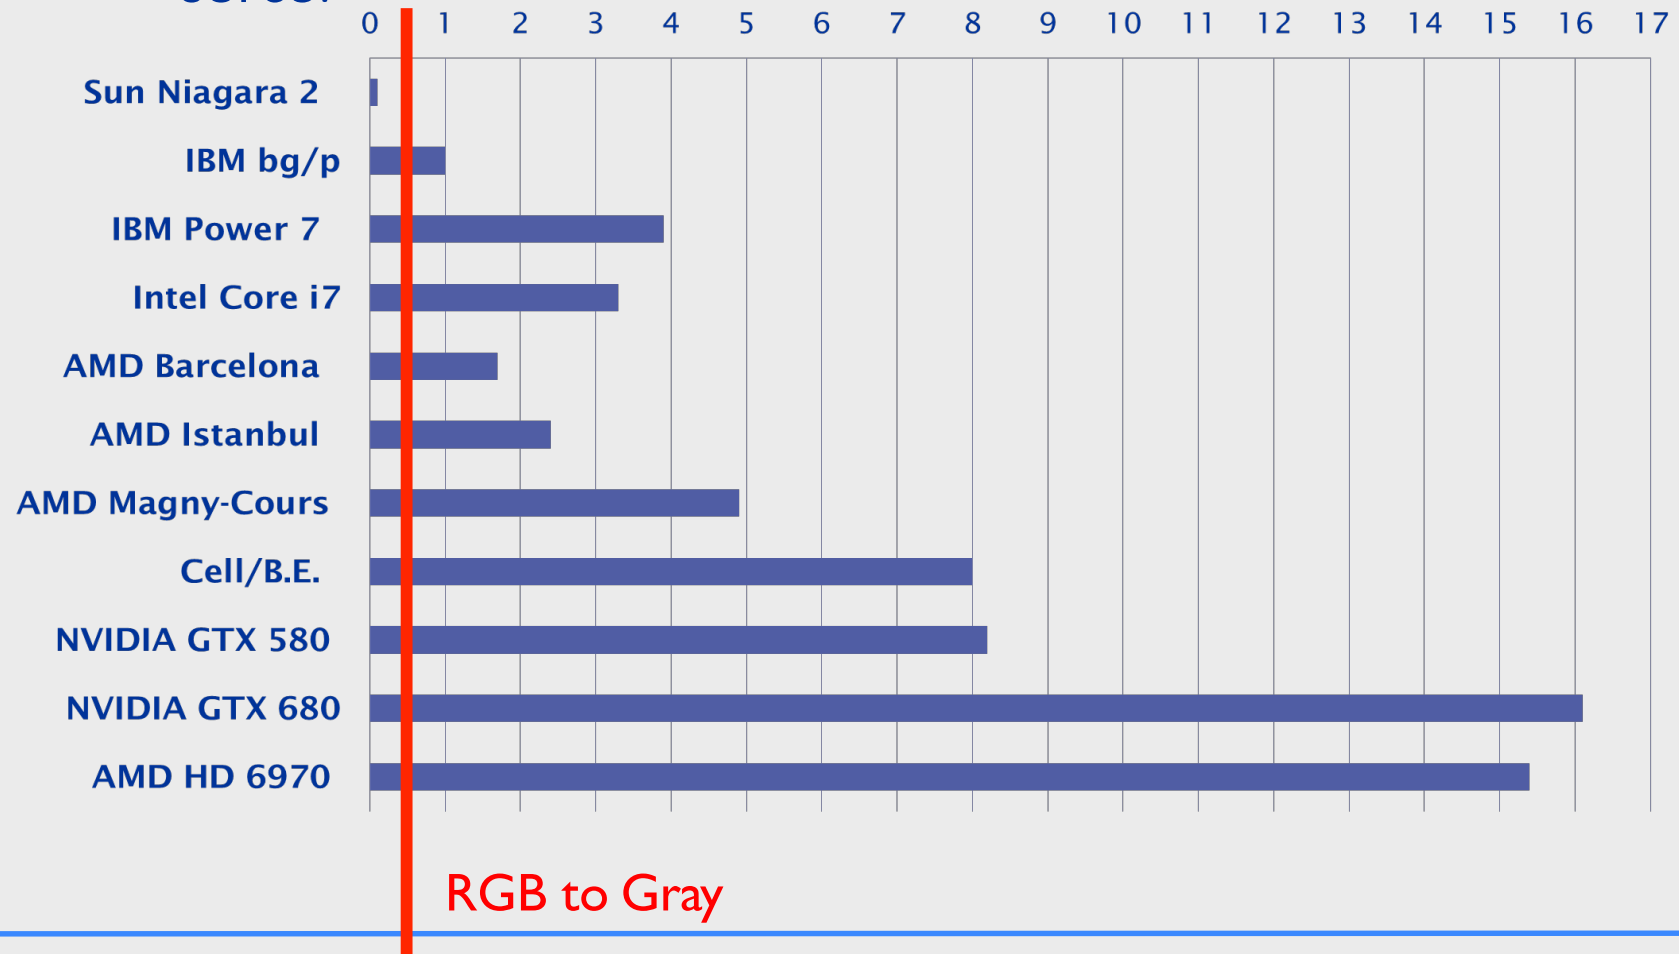
Arithmetic intensity :

```
for (int y = 0; y < height; y++) {  
    for (int x = 0; x < width; x++) {  
        Pixel pixel = RGB[y][x];  
        gray[y][x] =  
            0.30 * pixel.R  
            + 0.59 * pixel.G  
            + 0.11 * pixel.B;  
    }  
}
```

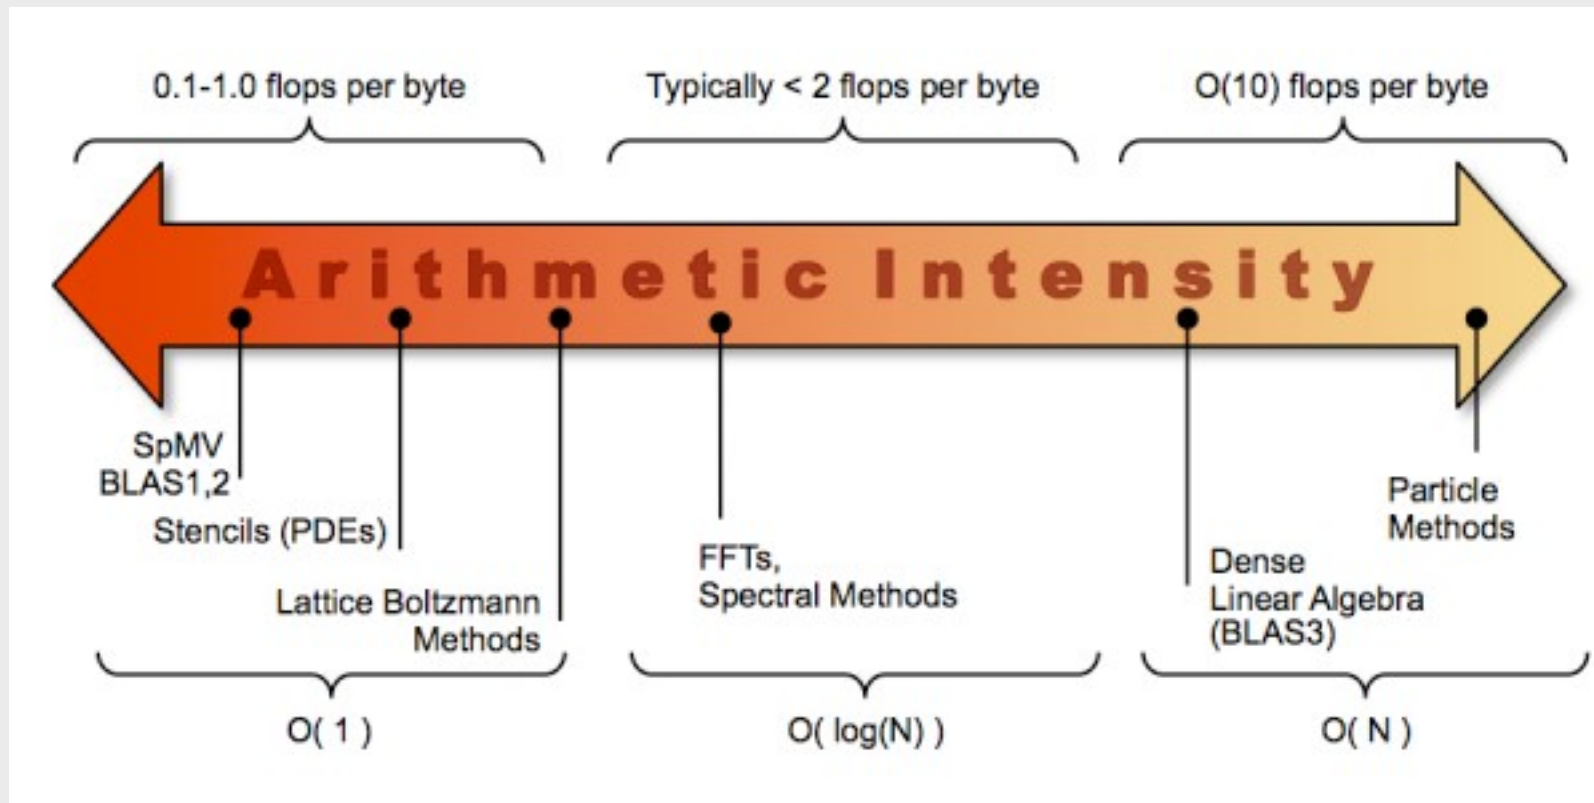
- What is the arithmetic intensity of the RGB-gray kernel?
- Is this a compute-intensive or a memory-intensive kernel?
- Is the kernel type (memory- or compute-intensive) dependent on the application, machine, or both?

Compute or memory intensive?

Arithmetic intensity for several actual many-cores:



AI for classes of applications



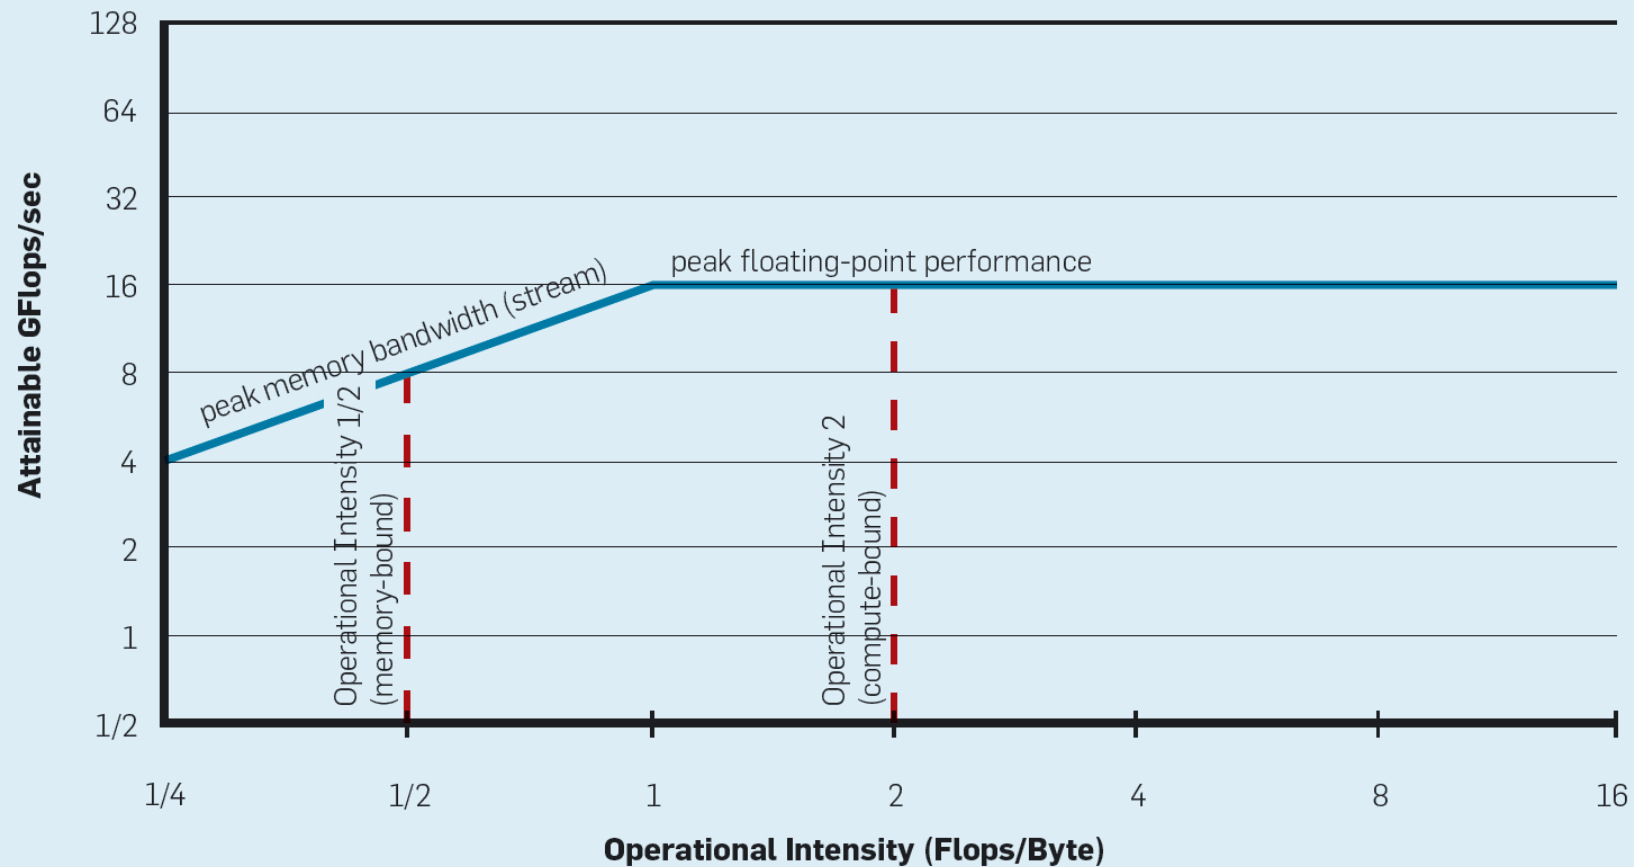
Operational intensity

- An extension (i.e., refinement) of arithmetic intensity:
The **number of operations** per **byte of memory** accessed
- Difference with Arithmetic Intensity
 - operations, not just arithmetic
 - caches are “skipped”
 - *count operations only*
 - *transfers not between processor and cache*
 - *but between cache and DRAM memory*

Attainable performance

- Attainable GFlops/sec
= $\min (\text{Peak Floating-Point Performance}, \text{Peak Memory Bandwidth} * \text{Operational Intensity})$
- To translate:
 - if an application is **compute-bound** => performance is limited by peak performance
 - if an application is **memory-bound** => performance is limited by the load it puts on the memory system

The Roofline model



AMD Opteron X2 (two cores): 17.6 gflops, 15 GB/s, ops/byte = 1.17

Attainable performance (cont'd)

- It is a measure of the performance an application **can** achieve
 - *more realistic because of the AI*
- Typical case: application A runs on platform X in $T_{\text{exec_A}}$:

```
PeakCompute(X) = maxFLOP GLOPS/s           (catalogue)
PeakBW(X) = maxBW GB/s                       (catalogue)
RooflineCompute(A,X) = min(AI(A)*maxBW, maxFLOP) (model)
AchievedCompute(A,X) = FLOPs(A)/Texec_A      (real execution)
AchievedBW(A,X) = MemOPs(A)/Texec_A          (real execution)

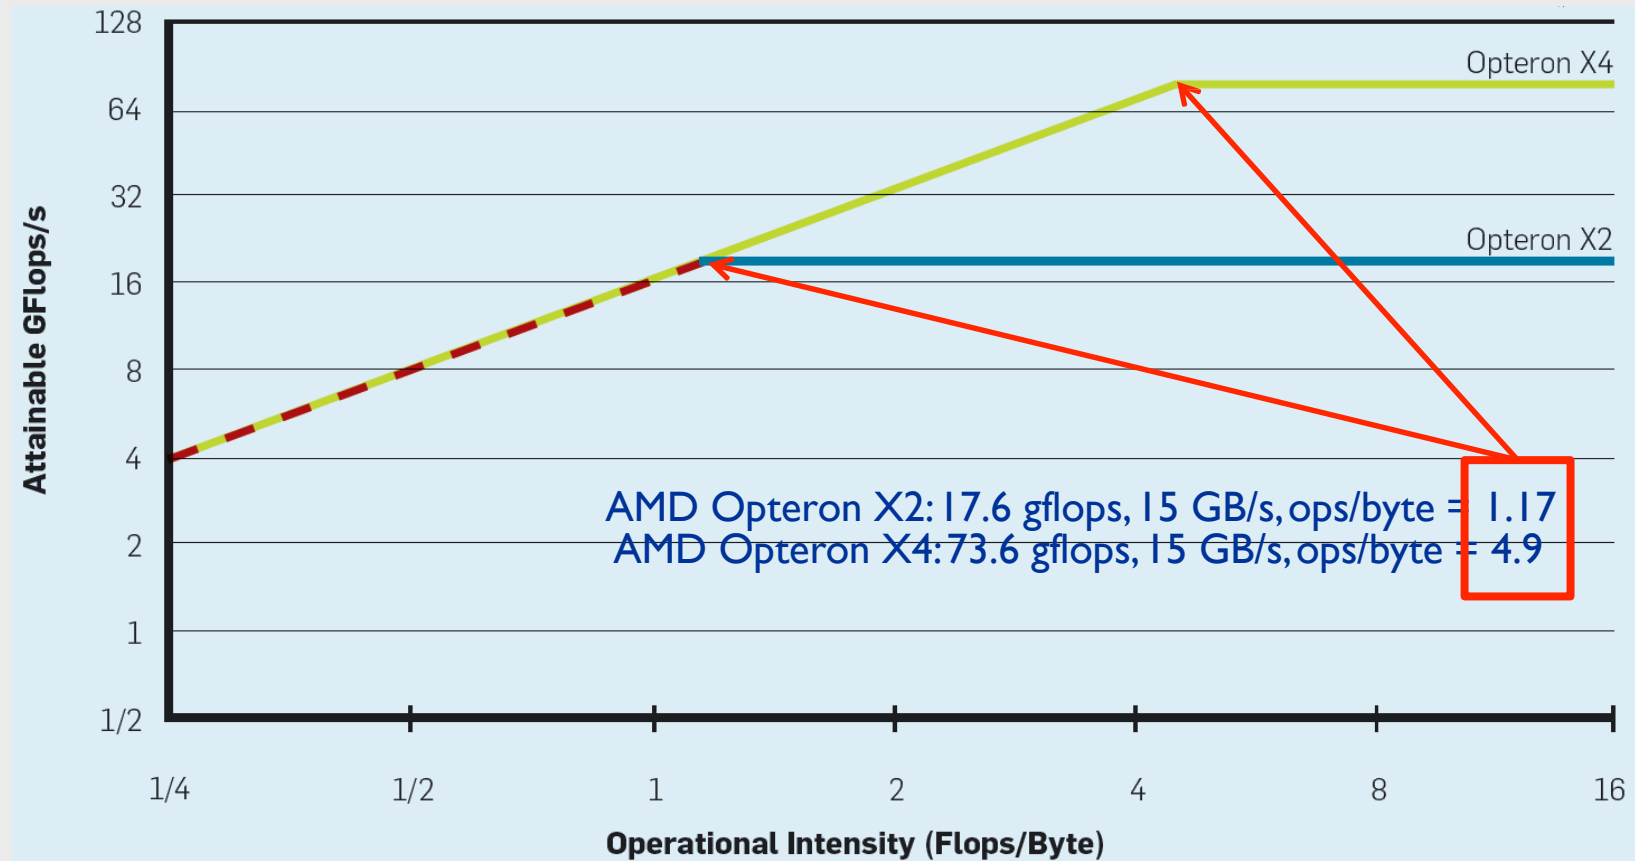
UtilizationCompute = AchievedCompute(A,X)/PeakCompute(X) < 1
UtilizationBW = AchievedBW(A,X)/PeakBW(X) < 1
```

PeakCompute >= **Roofline** > **AchievedCompute**

PeakBW ? **AchievedBW**

Note: AchievedBW > PeakBW <=> faster memories on the chip play a role.

Roofline: comparing architectures



Use the Roofline model

Determine what to do first to gain performance

- increase memory streaming rate (fights mem-boundness)
 - *GPU: memory coalescing*
 - *CPUs: better caching*
- apply in-core optimizations (fights compute-boundness)
 - *vectorization*
- increase arithmetic intensity (fights mem-boundness)
 - *change your algorithm*
 - *think of new ways to reuse the data*