

Asynchronous Parallel SGD

Reference:

J. Keuper and F-J. Pfreundt, Asynchronous Parallel Stochastic Gradient Descent:
A numeric core for scalable distributed machine learning algorithms, in Proc. MLHP 2015.

1. Gradient Descent Optimization

Batch Optimization: Update ‘weights’ w_i by mean gradient generated by ALL samples.

Algorithm 1 BATCH optimization with samples $X = \{x_0, \dots, x_m\}$, iterations T and states w

```
1: for all  $t = 0 \dots T$  do  
2:   Init  $w_{t+1} = 0$   $w_t$   
3:   for all  $x_j \in X$  do  
4:     aggregate  $w_{t+1} = w_{t+1} + \partial_w x_j(w_t)$   
5:    $w_{t+1} = w_{t+1} / |X|$ 
```

Partial derivative w.r.t. x_j : $\Delta_j(w_t) := \partial_w x_j(w_t)$

Stochastic Gradient Descent (SGD):

Update 'weights' w_t for every single sample. (stochastic: select a training sample at random)

Algorithm 2 SGD with samples $X = \{x_0, \dots, x_m\}$, iterations T , steps size ϵ and states w

Require: $\epsilon > 0$

- 1: **for all** $t = 0 \dots T$ **do**
 - 2: **draw** $j \in \{1 \dots m\}$ uniformly at random
 - 3: **update** $w_{t+1} \leftarrow w_t - \epsilon \partial_w x_j(w_t)$
 - 4: **return** w_{T+1}
-

Mini-batch SGD:

Update weights w_t after a group of training samples (mini-batch).

Algorithm 4 Mini-Batch SGD with samples $X = \{x_0, \dots, x_m\}$, iterations T , steps size ϵ , number of threads n and mini-batch size b

Require: $\epsilon > 0$

- 1: **for all** $t = 0 \dots T$ **do**
 - 2: **draw** mini-batch $M \leftarrow b$ samples from X
 - 3: **Init** $\Delta w_t = 0$
 - 4: **for all** $x \in M$ **do**
 - 5: **aggregate update** $\Delta w \leftarrow \partial_w x_j(w_t)$
 - 6: **update** $w_{t+1} \leftarrow w_t - \epsilon \Delta w_t$
 - 7: **return** w_{T+1}
-

2. Parallel SGD

n nodes/threads operate independently until convergence.

Algorithm 3 SimuParallelSGD with samples $X = \{x_0, \dots, x_m\}$, iterations T , steps size ϵ , number of nodes n and states w

Require: $\epsilon > 0, n > 1$

```
1: define  $H = \lfloor \frac{m}{n} \rfloor$ 
2: randomly partition  $X$ , giving  $H$  samples to each node
3: for all  $i \in \{1, \dots, n\}$  parallel do
4:   randomly shuffle samples on node  $i$ 
5:   init  $w_0^i = 0$ 
6:   for all  $t = 0 \dots T$  do
7:     get the  $t$ th sample on the  $i$ th node
8:     update  $w_{t+1}^i \leftarrow w_t^i - \epsilon \Delta_t(w_t^i)$ 
9: aggregate  $v = \frac{1}{n} \sum_{i=1}^n w_t^i$ 
10: return  $v$ 
```

M. Zinkevich, et al, Parallelized stochastic gradient descent. In Proc. NIPS 2010, pp. 2595-2603.

Asynchronous SGD

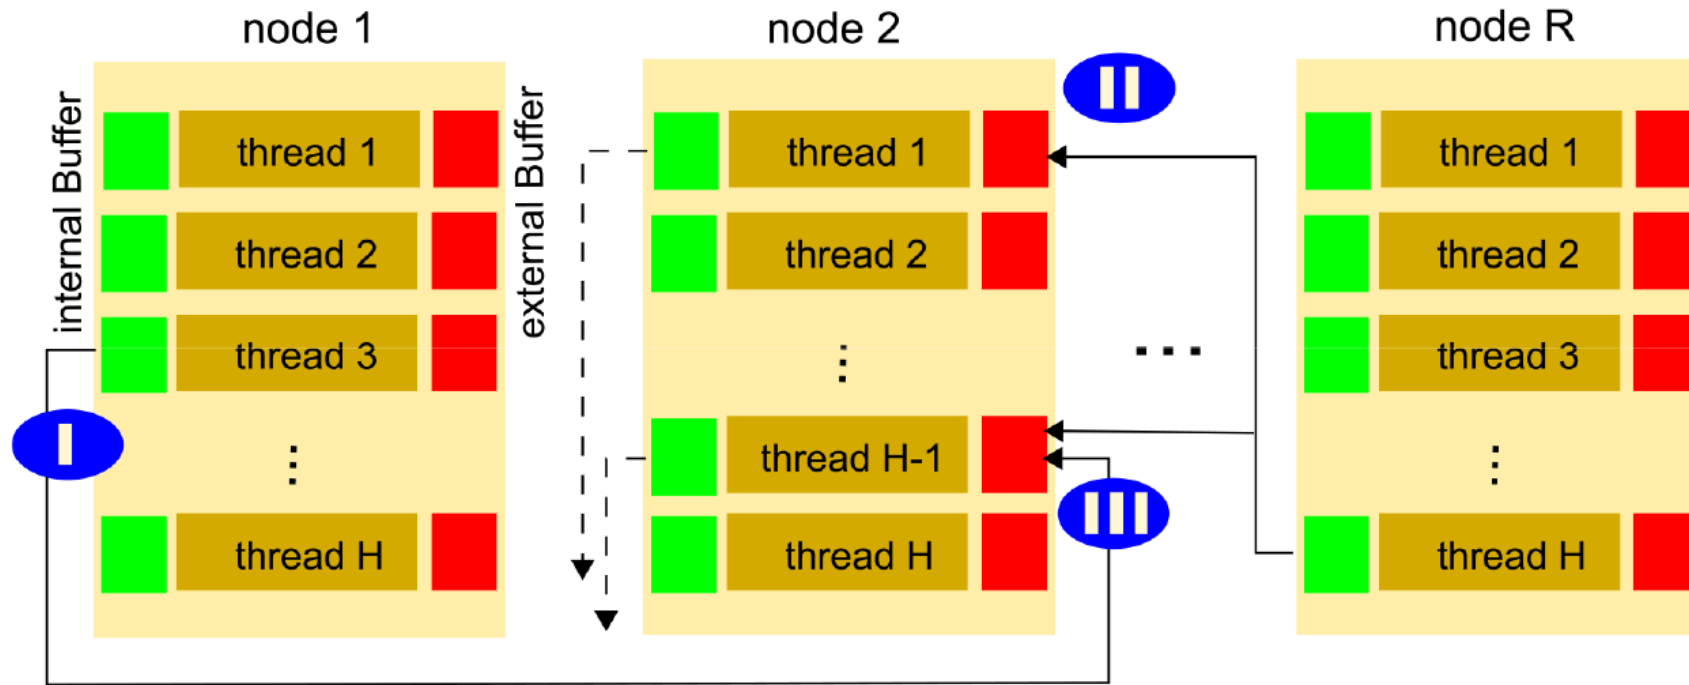


Figure 2: Overview of the asynchronous update communication used in ASGD. Given a cluster environment of R nodes with H threads each, the blue markers indicate different stages and scenarios of the communication mode. I: Thread 3 of node 1 finished the computation of its local mini-batch update. The external buffer is empty. Hence it executes the update locally and sends the resulting state to a few random recipients. II: Thread 1 of node 2 receives an update. When its local mini-batch update is ready, it will use the external buffer to correct its local update and then follow I. III: Shows a potential data race: two external updates might overlap in the external buffer of thread $H-1$ of node 2. Resolving data races is discussed in section 4.4.

3. ASGD

Algorithm 5 ASGD ($X = \{x_0, \dots, x_m\}, T, \epsilon, w_0, b$)

Require: $\epsilon > 0, n > 1$

- 1: **define** $H = \lfloor \frac{m}{n} \rfloor$
 - 2: randomly **partition** X , giving H samples to each node
 - 3: **for all** $i \in \{1, \dots, n\}$ **parallel do**
 - 4: randomly **shuffle** samples on node i
 - 5: **init** $w_0^i = 0$
 - 6: **for all** $t = 0 \dots T$ **do**
 - 7: **draw** mini-batch $M \leftarrow \underline{b \text{ samples from } X}$
 - 8: **update** $w_{t+1}^i \leftarrow w_t^i - \epsilon \Delta_M(w_{t+1}^i)$
 - 9: **send** w_{t+1}^i to random node $\neq i$
 - 10: **return** w_I^1
-

ASGD updating: handling of multiple updates in one iteration at node i

To combine with the communicated state $w_{t'}^j$ from an unknown iteration t' of some random node j :

$$\overline{\Delta_t(w_{t+1}^i)} = \frac{1}{2} (w_t^i + w_{t'}^j) + \Delta_t(w_{t+1}^i) \quad (2)$$

For the usage of N external buffers per thread, we generalize equation (2) to:

$$\overline{\Delta_t(w_{t+1}^i)} = w_t^i - \frac{1}{|N|+1} \left(\sum_{n=1}^N (w_{t'}^n) + w_t^i \right) + \Delta_t(w_{t+1}^i),$$

$$\text{where } |N| := \sum_{n=0}^N \lambda(w_{t'}^n), \quad \lambda(w_{t'}^n) = \begin{cases} 1 & \text{if } \|w_{t'}^n\|_2 > 0 \\ 0 & \text{otherwise} \end{cases}$$

Parzen-window: to include only the “good” update from other threads:

$$\delta(i, j) := \begin{cases} 1 & \text{if } \|(w_t^i - \epsilon \Delta w_t^i) - w_{t'}^j\|^2 < \|w_t^i - w_{t'}^j\|^2 \\ 0 & \text{otherwise} \end{cases}$$

ASGD updating: illustration

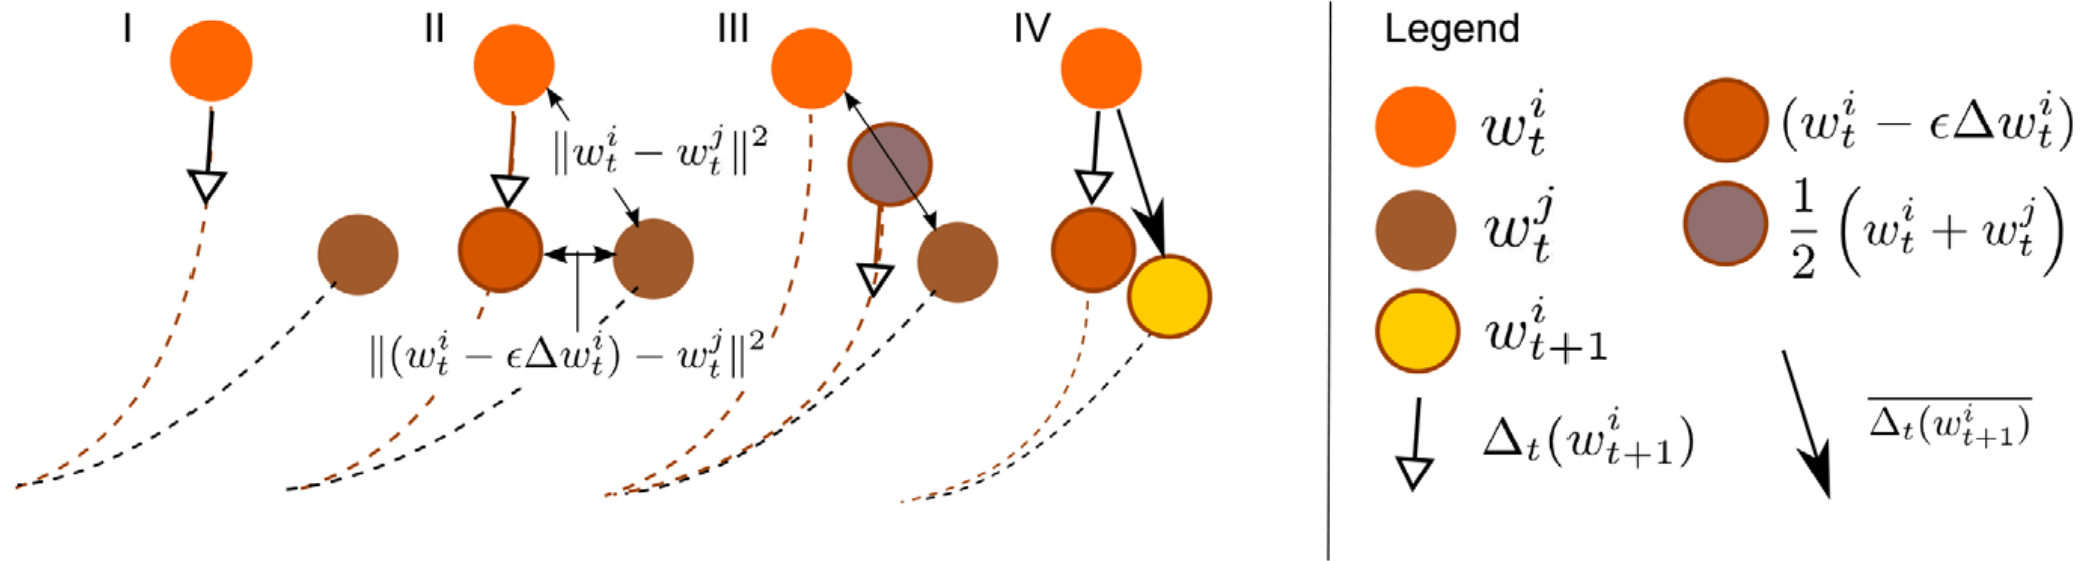
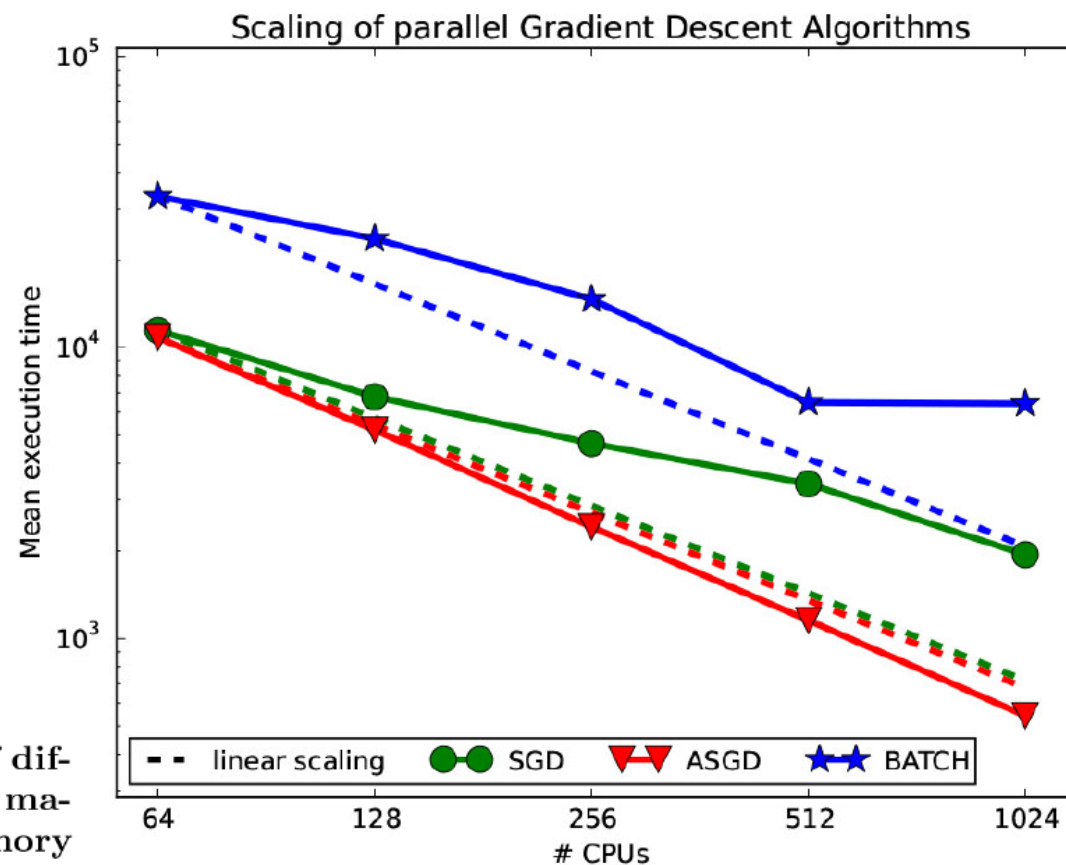


Figure 4: ASGD updating. This figure visualizes the update algorithm of a process with state w_t^i , its local mini-batch update $\Delta_t(w_{t+1}^i)$ and received external state w_t^j for a simplified 1-dimensional optimization problem. The dotted lines indicate a projection of the expected descent path to an (local) optimum. **I:** Initial setting: $\Delta_M(w_{t+1}^i)$ is computed and w_t^j is in the external buffer. **II:** Parzen-window masking of w_t^i . Only if the condition of equation (4) is met, w_t^j will contribute to the local update. **III:** Computing $\overline{\Delta_M(w_{t+1}^i)}$. **IV:** Updating $w_{t+1}^i \leftarrow w_t^i - \epsilon \overline{\Delta_M(w_{t+1}^i)}$.

4. Performance

Figure 1: Evaluation of the scaling properties of different parallel gradient descent algorithms for machine learning applications on distributed memory systems. Results show a K-Means clustering with $k=10$ on a 10-dimensional target space, represented by $\sim 1\text{TB}$ of training samples. Our novel ASGD method is not only the fastest algorithm in this test, it also shows better than linear scaling performance. Outperforming the SGD parallelization by [19] and the MapReduce based BATCH [5] optimization, which both suffer from communication overheads.



Asynchronous communication

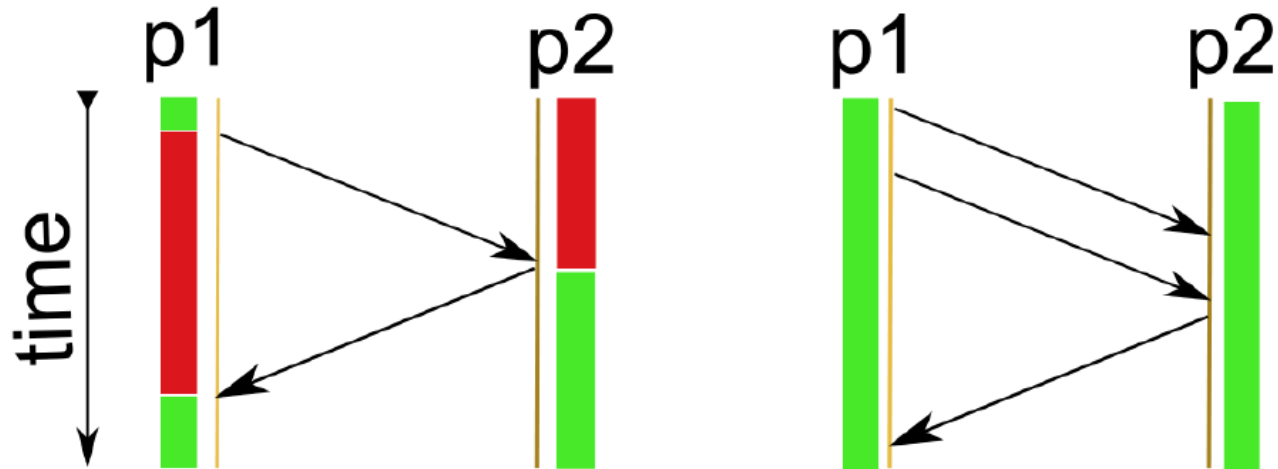


Figure 3: Single-sided asynchronous communication model (right) compared to a typical synchronous model (left). The red areas indicate dependency locks of the processes $p1, p2$, waiting for data or acknowledgements. The asynchronous model is lock-free, but comes at the price that processes never know if and when and in what order messages reach the receiver. Hence, a process can only be informed about past states of a remote computation, never about the current status.

Remark:
Implementation: one-sided
put and get, or asyn.
communication.

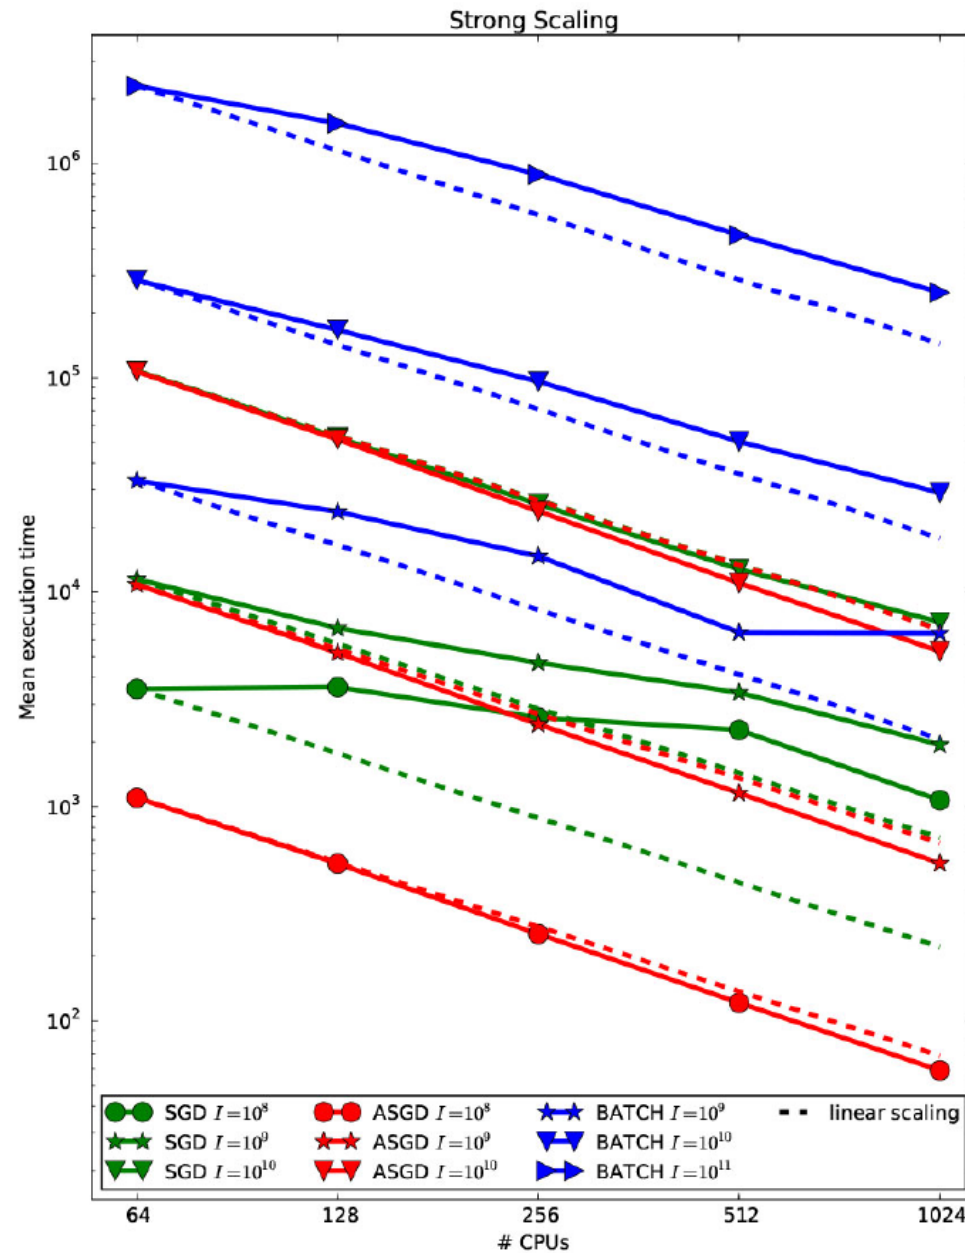


Figure 5: Results of a strong scaling experiment on the synthetic dataset with $k=10$, $d=10$ and $\sim 1\text{TB}$ data samples for different numbers of iterations I . The related error rates are shown in figure 9.

Convergence Speed

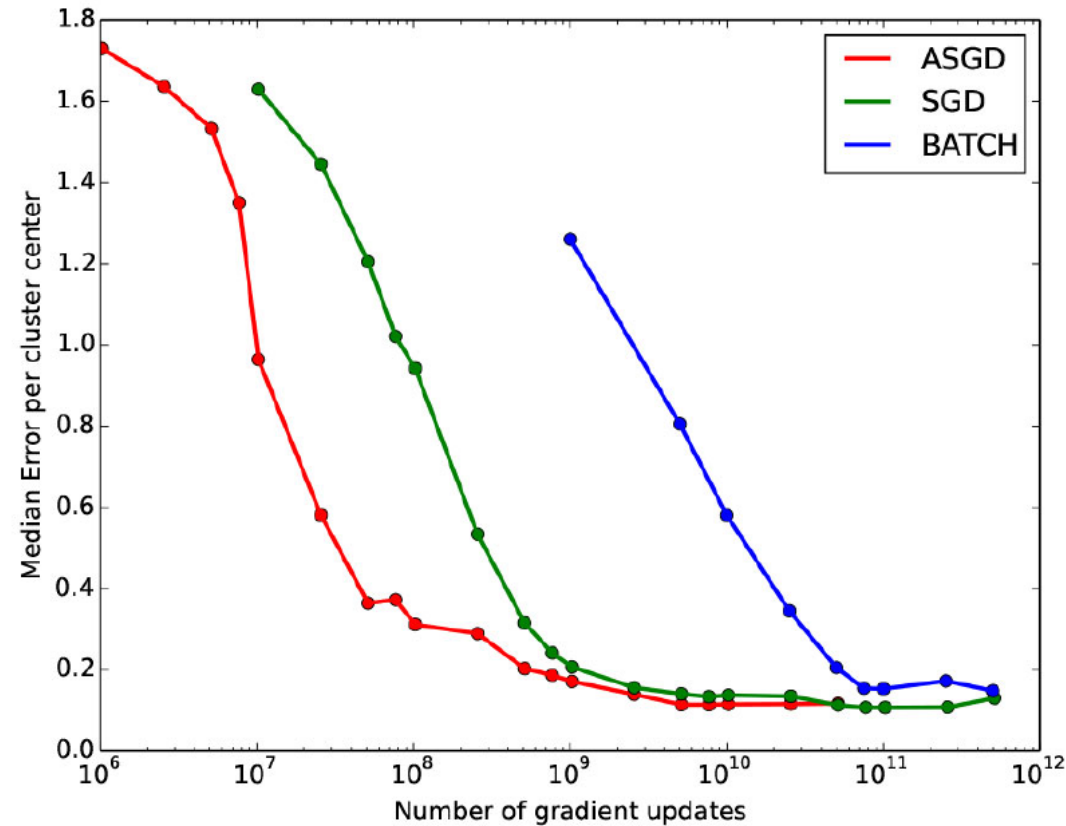


Figure 8: Convergence speed of different gradient descent methods used to solve K-Means clustering with $k = 100$ and $b = 500$ on a 10-dimensional target space parallelized over 1024 CPUs on a cluster. Our novel ASGD method outperforms communication free SGD [19] and MapReduce based BATCH [5] optimization by the order of magnitudes.

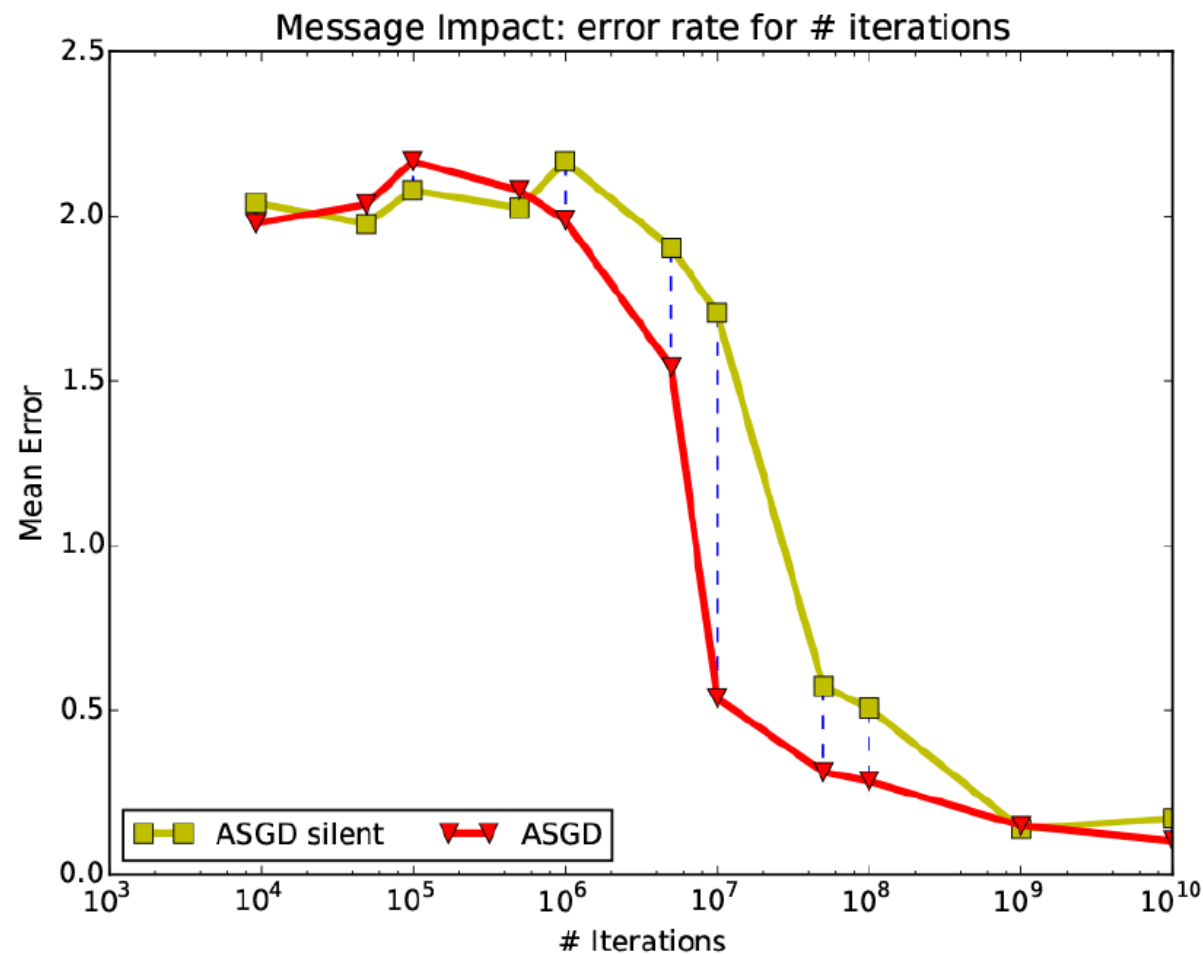


Figure 14: Convergence speed of ASGD optimization (synthetic dataset, $k = 10, d = 10$) with and without asynchronous communication (silent).